

When the Regime Returns: Retention and Invariance in Decision-Focused Strategy Pools for Non-Stationary Markets

Yuhao Li
University of Pennsylvania
li88@sas.upenn.edu

Abstract

Financial markets are non-stationary in two distinct ways that existing strategy-allocation methods conflate. **Between regimes**, market states (trend, range, high-volatility, crisis) switch and recur on irregular schedules with long dormancy, so a reward-chasing capital allocator—recent-P&L weighting or a learned mixture-of-experts router—starves and overwrites the dormant-regime specialist and pays a full recalibration cost exactly when the regime returns. **Within a regime**, each recurrence is a different draw from the regime’s environment distribution (the 2008 crisis is not 2020 is not 2022), so a specialist trained by empirical risk minimization on past episodes of its own regime still fails out-of-distribution on the next one. We show these failure modes are governed by *independent* design choices and propose RISP (Regime-Invariant Specialist Pools), a pool of capacity-bounded, decision-focused specialists in which (i) regimes are acquired by batched winner-take-all competition, yielding a converged regime-to-specialist assignment that is *independent of current reward*—the crisis specialist idles through calm markets and retains its calibration by identity, not P&L—and (ii) each specialist trains its niche with a decision-focused *invariance* objective (variance plus mean of an SPO-style regret surrogate across historical episodes of its regime). We prove a post-reactivation regret bound that decomposes into a *forgetting deficit*—zero for any isolated covering assignment, approaching the full relearning cost for any reward-driven one—and an *invariance gap* controlled by the episode-variance penalty, including the Pinsker-type bridge from divergence-metric invariance to regret. In controlled regime-switching markets with realistic signal-to-noise, the predicted 2×2 dissociation holds with all pre-registered orderings confirmed (100 seeds, Holm-corrected): emergent reward-independent assignment cuts post-reactivation decision regret by 4.4% versus a total-capacity-matched learned router ($p=2.5 \times 10^{-7}$; 4.3% versus the per-specialist-matched monolith, $p=1.9 \times 10^{-7}$, capacity accounting in §5.2), the invariance objective cuts it a further 5.0% ($p=5.8 \times 10^{-9}$; 9.1% jointly, $p=6.0 \times 10^{-23}$; 36% of the excess over the converged invariant-service floor), and the interaction is *super-additive*: invariant training buys nothing in a reward-driven monolith ($p=0.89$ vs. ERM), because an objective cannot help a head that has been evicted. The learned router and the recent-performance allocator inherit the monolith’s failure; the full system is statistically indistinguishable from a hand-pinned oracle assignment trained with the same objective ($p=0.53$)—recovering the skyline without being handed the assignment—and beats the oracle assignment trained with ERM ($p=1.5 \times 10^{-6}$). We pair the mechanism with an honest audit in the GAUSE tradition, organized as a *two-gate* deployment protocol whose gates are checkable before any mechanism is bought (gate 1: regime structure on the decision metric against shuffled labels; gate 2: a positive measured forgetting deficit, monolith minus oracle-pinned). Daily crypto and commodity panels fail gate 1 and dissociate null, as required; 4-hour crypto passes gate 1 but measures a zero deficit and is null again; on *drawdown-labeled daily equities* (49 industry portfolios, 1990–2025, pre-registered) the deficit is **positive on real data** ($+0.00089 \pm 0.00011$, 100 seeds; seed-noise interval, single history) and the full pre-registered ordering emerges (headline contrast Holm $p=5.2 \times 10^{-65}$; the ordering’s

binding pair at 1.6×10^{-15}). The record now extends beyond one sample: on a **withheld era** (36 industries, 1926–1989, frozen specification), the deficit’s *sign* predicted the arm-ordering direction in **5 of 6** pre-registered cells—in both directions, including two positive-deficit cells reproducing the full ordering out-of-sample (Holm 1.9×10^{-6} , 2.1×10^{-4})—while gate 1’s conditioning criterion tracked neither and is demoted to a structure screen. Transaction costs *amplify* rather than kill the effect (net-of-cost deficit $+0.00126 \rightarrow +0.00230$ as crisis-window costs sweep 25–100 bps; the one observed inversion, under the vol-band labeler, reverses net of costs); a decoupled-replay monolith closes only 56% of the deficit in the controlled market and, on equities, collects the deficit yet still loses to the emergent pool (Holm 0.013); and the deficit is not an artifact of the discrete memory model—a shared-trunk MLP with nothing evicted reproduces it as pure gradient interference ($+0.01009 \pm 0.00037$, 100 seeds). One attribution is honestly revised: on real data the deficit neither decays across the probe window nor follows any dormancy or rehearsal covariate—it behaves as a persistent *regime-conditional allocation deficit*, not demonstrated eviction-forgetting, which remains established only in the controlled settings. The first externally-lodged temporal test—an announcement-lagged NBER-recession labeler with no free threshold, its registration deposited with the Open Science Framework before the recession dates were ever joined to the panel—has since resolved as a hit under both scoring forms ($\hat{\Gamma} = +0.00081 \pm 0.00033$ with the ordering at Holm 1.6×10^{-4} , cost-amplified; its schedule-resampling counterpart flat, disclosed). Caveats carried at equal prominence: each walk-forward window is a single market history (four windows are now scored), the deficit’s threshold window relocates across eras, aggregation itself can manufacture slow dynamics, and the constituent-level CRSP test—its registration lodged with OSF before any data access—remains the decider. All claims are regret and retention claims, not alpha claims; code, data, and per-figure result files are released.

1 Introduction

A multi-strategy fund maintains playbooks for market conditions that are mostly absent: the crisis desk idles through years of calm; the short-volatility book hibernates through stress. The allocation question that governs such a system—*which strategy gets capital, attention, and recalibration today, and who keeps a strategy whose regime is dormant?*—has a standard machine-learning answer (route by recent reward: a learned gate, a recent-Sharpe weighting) and a standard failure mode that the GAUSE companion study isolated in controlled prediction domains (Li, 2026): a dormant regime emits no reward, so any allocator that protects capacity using current reward is structurally blind to exactly the knowledge that most needs protection. When the regime returns, the system relearns what it once knew, and in markets the relearning window—the first days of a crisis—is precisely when decisions are most consequential.

This paper develops that observation into a theory and system for *decision-focused* learning in non-stationary markets, and in doing so isolates a second, orthogonal failure mode that the retention story alone misses. Suppose the crisis specialist *is* protected—its parameters untouched through dormancy. Protection guarantees only that the specialist still knows *the last crisis*. Each occurrence of a regime is a distinct draw from that regime’s environment distribution: flight-to-quality correlations, volatility clustering, and liquidity spirals recur *structurally*, but the surface features that carried them in 2008 differ from 2020 and from 2022. A specialist trained by empirical risk minimization (ERM) on its own regime’s history loads on episode-specific ("spurious") regularities and fails out-of-distribution on the next episode (Arjovsky et al., 2019; Yan et al., 2026). Retention solves *between-regime* forgetting; it does nothing for *within-regime* drift.

Our central claim is that these are **two orthogonal kinds of non-stationarity**, addressed by *independent* mechanisms:

- **(N1) Between-regime switching** is an *allocation* problem. It is solved by any capacity assignment that *isolates* the dormant specialist—no foreign updates during its regime’s dormancy—while persistently covering the regime inventory, so that the specialist retains by identity; converged reward-independent competition is one way to obtain such isolation, and reward-chasing allocation destroys it. No training objective can substitute: a forgotten head is gone regardless of the loss it was trained with.

- **(N2) Within-regime drift** is a *training-objective* problem. It is solved by an invariance objective across the regime’s historical episodes—penalizing the variance of a decision-focused (SPO-style) regret surrogate across episodes so the specialist learns the decision-relevant structure common to *all* crises rather than the idiosyncrasies of the last one. No allocation mechanism can substitute: retention of an episode-overfit head retains the overfitting too.

The retention half of this division has a recent formal ancestor in the average-reward account of continual learning (Kumar et al., 2025), which argues that a computationally bounded agent should forget freely whatever will not recur—forgetting non-recurring information is not catastrophic—and that even recurring information need not be retained when relearning on arrival is cheap: their capacity-bounded agent facing a two-season market relearns each season’s function as it arrives. Our deficit $\hat{\Gamma}_{\text{forget}}$ (Contribution C4) is a decision-regret operationalization of exactly that trade-off for markets—a per-substrate measurement of whether the regime inventory recurs in a way a bounded learner is paying to relearn—and Proposition 4 prices their relearn-quickly proviso.

The paper’s single falsifiable claim is the resulting 2×2 : crossing allocation mechanism (reward-driven vs. reward-independent) with training objective (ERM vs. invariant) should produce monotone improvement along both axes in post-reactivation decision regret, with a signature asymmetry—the invariance axis should be inert in the reward-driven row, because an evicted head cannot benefit from the objective that trained it. Our experiments confirm every cell of this prediction, including the asymmetry (Section 5.2), and our theory derives it (Section 4).

Scope statement (the sleeping-experts deflation, addressed up front). The contribution concerns capacity-bounded *learners*, not combinations of fixed experts. Sleeping-experts and adaptive-regret schemes (Freund et al., 1997; Hazan and Seshadhri, 2009; Daniely et al., 2015) also freeze dormant experts’ state—but only because their experts are *static strategies* with nothing to forget. The problem RISP solves arises precisely when specialists must keep *learning* within their niche under bounded capacity, so that reward-driven reallocation causes destructive interference across regimes. We operationalize this distinction as a baseline pair: Hedge over *fixed* strategies retains trivially but cannot adapt within-regime and is dominated by every learning arm (Section 5.2); Hedge over capacity-bounded *learning* experts inherits the interference failure and tracks the monolith. The most dangerous related work thus becomes a supporting experiment.

Contributions.

1. **Problem decomposition (C1).** We formally decompose non-stationary financial decision error into a between-regime *forgetting deficit* and a within-regime *invariance gap*, and show—in theory and experiment—that they are governed by independent design choices (allocation vs. training objective), with a predicted and confirmed super-additive interaction: invariant training pays *only* inside a retaining allocation.
2. **Method (C2).** RISP: the first strategy-pool architecture combining reward-independent, emergent (competition-driven) capacity assignment with decision-focused invariant learning per niche. The anticipated financial adaptation of the GAUSE competition mechanism is a *batched* winner-take-all window (W_c days of cumulative within-regime regret); we pre-registered the expectation that per-step competition would converge to noise at realistic signal-to-noise, and report honestly that the expectation is *not* borne out—the exponentiated-gradient affinity accumulation across windows performs the necessary averaging, and the mechanism is robust across $W_c \in \{1, \dots, 60\}$ (Section 5.11). No gate is trained, no freezing schedule is set, no regime-to-desk map is designed.
3. **Theory (C3).** A post-reactivation regret bound at fixed tail level, $\mathbb{E}_A[\text{CVaR}_\delta[\rho_{\text{react}}]] \leq G_{\text{inv}}(\delta) + G_{\text{forget}}$ (Proposition 2)—allocation controls the mean of the reactivation transient, the invariance objective its tail across episode draws, a tail reading whose pre-registered quantile test subsequently *failed*, demoting it to an unconfirmed bound (see after Proposition 2)—in which $G_{\text{forget}} = 0$ for *any* isolated covering assignment while reward-driven assignment drives it to the full relearning cost at an explicit rate (Proposition 1); $G_{\text{inv}}(\delta)$ is controlled by the episode-variance penalty (Lemma 1), with an explicit Pinsker-type bridge from divergence-metric invariance statements to the regret metric (Lemma 2)—closing the

technical gap between Inv-PnCO-style guarantees (Yan et al., 2026) and decision regret. A break-even analysis (Proposition 4) gives the dormancy window in which retention pays once idle-specialist carrying costs are charged.

4. **Evaluation protocol (C4).** A diagnostic-first standard for regime-aware financial ML: *post-reactivation regret* in a probe window after each reactivation, leakage-free causal regime labelings, and a pre-deployment protocol run *before* claiming gains — a structure *screen* plus a binding diagnostic. We state the protocol as its own record left it, not as designed: the screen’s outcomes have not tracked pool value in any tested cell (three dissociations, including its strongest pass coinciding with an inversion), so it screens for learnable structure and licenses nothing; the diagnostic is the measured deficit. *Screen* (structure): regime-conditioning must beat block-shuffled labels on the decision metric. *Gate 2* (harvestable deficit): the measured deficit—monolith minus oracle-pinned post-reactivation regret—must be positive. The decomposition of Proposition 2 bounds this quantity’s *transient* branch (eviction-relearning), which is the measured mechanism in our controlled settings; on the real substrate where the probe profile shows no decay (Section 5.7), $\hat{\Gamma}$ measures a persistent quantity no proposition in this paper bounds—we say so rather than borrow the theorem’s authority. The record across the real substrates (Sections 5.4–5.7): daily crypto/commodities fail gate 1 and dissociate null, as required; 4-hour crypto passes gate 1, fails gate 2, dissociates null; on daily equities (1990–2025) the deficit’s *sign* predicted the walk-forward outcome in both tests—zero deficit $\Rightarrow$ the monolith wins significantly (the invariance premium with no claim), positive deficit $\Rightarrow$ the full pre-registered ordering emerges—and on a withheld 64-year era (1926–1989, frozen specification) the same sign rule scored 5 of 6 pre-registered cells, in both directions. The protocol was revised by its own record—gate 2 was added after one refutation, and gate 1, whose strongest pass in the program ($z=23.5$, 1926–1989) coincides with an *inverted* pool ordering, is demoted to a structure screen: it certifies that regime labels carry decision-relevant information, and it does not price the pool. What survives is a single-number deployment rule, $\hat{\Gamma}_{\text{forget}}$ from the A1-vs-A9 counterfactual—now tested out-of-sample in both directions and robust to measured crisis-window transaction costs—whose forward test on CRSP constituents is pre-committed. Where the deficit is not positive, no retention mechanism (ours included) can be bought, and a method that had claimed otherwise would be fitting noise.
5. **Honest audit (C5).** Conditions under which RISP does *not* pay, each quantified: slack capacity ($K \rightarrow R$ closes the monolith gap); high within-regime signal-to-noise (a fresh learner recalibrates within days, making retention worthless—we sweep the break-even relearning speed); short dormancy (the retention gap grows with dormancy length); structureless regimes (the E0 gate); and carrying costs above the break-even threshold. And where the deficit *is* positive on real data, we do not claim its mechanism: the measured quantity behaves as a regime-conditional *allocation* deficit (persistent, not decaying; Section 5.7), and eviction-forgetting proper is demonstrated only in the controlled settings (E1 and Section 5.3).

We emphasize register throughout: our claims are *regret and retention* claims about a mechanism, demonstrated where its preconditions verifiably hold, not alpha claims about live markets. The honest summary of the real data is that the two preconditions are *separable*, fail separately, and—on the substrates where both the structure and a positive measured forgetting deficit coexist (drawdown-labeled daily industries, Section 5.7)—jointly deliver the pre-registered ordering in single-history walk-forwards, at aggregate (portfolio) level, an effect that survives and strengthens under measured crisis-window transaction costs: no decision-metric structure on small daily panels (gate 1), no harvestable forgetting on the intraday substrate where structure exists (gate 2), an inversion where the deficit is zero or negative, the dissociation where it is positive, and a 5-of-6 out-of-sample scorecard for that sign rule on a withheld 64-year era—each finding surfaced by the diagnostic-first protocol, the first consistent with the GAUSE oracle audit on the prediction metric (Li, 2026).

2 Related Work

Regime-switching finance. Markov-switching models (Hamilton, 1989), regime effects in asset returns (Ang and Bekaert, 2002; Ang and Timmermann, 2012), regime-conditioned allocation (Guidolin and Timmer-

mann, 2007; Nystrup et al., 2018), and statistical jump models for persistent market states (Nystrup et al., 2020) detect regimes and condition a *single* model on them. None address bounded capacity or dormancy-driven forgetting of the specialist itself: the implicit assumption is that conditioning information is free to retain. Regime-switching forecast combination (Elliott and Timmermann, 2005) conditions the combination *weights* on an inferred regime but likewise assumes every component model remains fully estimated while its regime is dormant: the regime moves the weights, never the fitness of the experts; the regime-conditional allocation deficit of Section 5.7 measures precisely that assumption’s failure under bounded capacity. Our monolith arm (regime-conditioned, capacity-bounded, always trained on the active regime) is this literature’s best case under a capacity budget, and the comparison isolates what conditioning alone cannot provide.

Decision-focused learning (predict-then-optimize). The SPO framework and its convex surrogate (Elmachtoub and Grigas, 2022), surrogate relaxations for hard combinatorial problems (Mandi et al., 2020), learned decision losses (Shah et al., 2022), and the field survey (Mandi et al., 2024) all assume a stationary (or one-shot OOD) environment. Inv-PnCO (Yan et al., 2026) introduces invariant learning for predict-and-optimize under distribution shift, bounding the divergence between solution distributions across environments. We adapt its environment-indexed invariance to *recurring* regimes with dormancy—episodes of a regime as environments—and contribute the bridge from divergence-metric guarantees to regret (Lemma 2), plus the orthogonal allocation axis that the stationary setting never encounters. Online decision-focused learning has since acquired its first regret guarantees under time-varying distributions (Capitaine et al., 2026); that analysis treats drift as unstructured, whereas our setting’s difficulty is *recurrence with dormancy*—the property that makes retention a design axis at all—which drift-only regret bounds do not touch. Closest in application is ReCAP (Pan et al., 2026), which builds the retention machinery end to end for portfolio management: change-point detection on market-state features segments history into variable-length regimes, each distilled into a policy-difference vector stored in a library that is frozen at deployment (only a state-driven softmax gate and the current regime’s vector keep training), so that a returning regime is served by re-weighting stored vectors. This is parameter-isolation retention realized by design—detected regime identity playing the role task identity plays in the PackNet family above—evaluated by the standard averaged forgetting/transfer metrics rather than post-reactivation decision regret, with no retention counterfactual. ReCAP therefore *presumes* the question our protocol makes primary: it deploys retention unconditionally, with no pre-deployment test of whether the substrate offers a harvestable deficit. The diagnostic of Section 5.7 prices exactly such machinery before it is bought—and our record contains real substrates where the sign is negative and an always-on retention stack pays the premium with no claim to collect.

Online learning and ensembles. Hedge/multiplicative weights (Freund and Schapire, 1997; Arora et al., 2012), universal portfolios (Cover, 1991), online portfolio selection (Li and Hoi, 2014), and the sleeping-experts / adaptive-regret line (Freund et al., 1997; Hazan and Seshadhri, 2009; Daniely et al., 2015) are the closest formal relatives. The scope statement of Section 1 is our differentiation: these schemes combine *fixed* experts, for which dormancy is harmless by construction. When the experts are capacity-bounded *learners*, the combination scheme’s weights are reward-driven and the training resource follows the weights—reintroducing exactly the interference that retention requires avoiding. The tracking line sharpens the point: Fixed Share (Herbster and Warmuth, 1998) and especially mixing past posteriors (Bousquet and Warmuth, 2002), which attains logarithmic regret against a small set of *returning* experts (see also time-selection regret (Blum and Mansour, 2007)), solve the recurrence problem for the combiner’s *weights*—a returning expert is re-elevated almost instantly. They cannot solve it for the expert’s *parameters*: under bounded capacity the training resource follows the weights, so the dormant specialist that mixing past posteriors would re-elevate has meanwhile been overwritten—MPP retains the pointer, not the playbook, and our A8b arm equipped with Fixed-Share weighting would inherit the same interference. Our A8a/A8b baseline pair tests both readings directly.

Mixture-of-experts and continual learning. MoE theory for continual learning (Li et al., 2025) shows isolation mitigates forgetting *provided* experts are plentiful and the gate’s updates are eventually terminated. Recent routing-stabilization work reduces router-induced expert churn during *active* training (CP-MoE (Liu et al., 2026); StaR-MoE (Guo et al., 2026)), but leaves the dormancy problem untouched—a regime that emits no reward for hundreds of steps gives a stabilized router nothing to stabilize toward—and neither evaluates

post-reactivation regret, the metric on which the failure appears. GAUSE (Li, 2026) showed empirically that in the excluded regime (bounded capacity, gate still learning) a reward-driven router inherits the monolith’s forgetting, and that emergent competitive exclusion reaches the required reward-independent assignment with no freezing schedule. We import that mechanism as a premise, port it to decision-focused financial learning (batched competition, regret-based quality), and extend the theory from prediction error to decision regret. Among classical continual-learning remedies, regularization approaches (EWC (Kirkpatrick et al., 2017); survey (Parisi et al., 2019)) operate within a single model, but the parameter-isolation family—PackNet (Mallya and Lazebnik, 2018), HAT (Serrà et al., 2018), SupSup (Wortsman et al., 2020), Expert Gate (Aljundi et al., 2017), PathNet (Fernando et al., 2017), progressive nets (Rusu et al., 2016)—already realizes reward-independent capacity assignment when task identity is given (or, for SupSup, inferred at test time), and is the continual-learning instantiation of the isolation property of Proposition 1. What that family does not provide is the allocation-under-competition analysis—which reward-driven assignment destroys isolation, at what rate, and how a covering assignment emerges without a designed task-to-mask map—nor the within-regime invariance axis, which is orthogonal to isolation and is where the super-additive interaction lives. Continual learning has also been proposed for investment decision-making directly (Philps et al., 2018), without either analysis.

Multi-strategy practice and strategy decay. The practitioner analogue of our mechanism is the siloed desk structure of multi-strategy funds; the practitioner analogue of our failure mode is recent-P&L capital allocation *across* desks. Strategy decay (McLean and Pontiff, 2016) and backtest-overfitting critiques (Bailey and López de Prado, 2014; Harvey and Liu, 2016; López de Prado, 2018) motivate both our within-regime drift axis and our statistical protocol (deflated statistics, no test-period tuning, controlled semi-synthetic ground truth).

3 Problem Setup and Method

3.1 Data stream, episodes, and the decision layer

A stream $t = 1, \dots, T$ presents observable per-asset features $X_t \in \mathbb{R}^{n \times d}$ and an unknown coefficient vector $y_t \in \mathbb{R}^n$ (next-period returns, $\|y_t\|_\infty \leq y_{\max}$), with a latent regime $r_t \in \{1, \dots, R\}$ evolving in persistent blocks. Regimes recur: the e -th maximal active spell of regime r is *episode* (r, e) , a stationary environment with law $P_{r,e}$ over (x, y) pairs; the dormancy $D(r)$ of a regime is the length of the inactive spell between consecutive episodes. Crisis-type regimes have long, irregular dormancy and few historical episodes—the regime that most needs a calibrated specialist is the one observed least and least recently.

Definition 1 (Decision layer and regret). *The feasible set is the cardinality-constrained long-only portfolio $Z = \{z \in [0, w_{\max}]^n : \sum_i z_i \leq W, \|z\|_0 \leq k\}$ with linear utility $F(z, y) = y^\top z$. The plug-in policy decides $z(\hat{y}) = \arg \max_{z \in Z} \hat{y}^\top z$ (exactly solvable: $w_{\max}$ on the k largest positive coordinates when $W \geq k w_{\max}$), and the per-step decision regret is*

$$\rho(\hat{y}; y) = F(z^*(y), y) - F(z(\hat{y}), y) \in [0, B], \quad B = 2k w_{\max} y_{\max}.$$

Regret—not prediction error—is the object of interest: a predictor can have large mean-squared error and zero regret (errors that do not change the top- k ranking are free) and small error but large regret (errors concentrated at the selection boundary). This is the core predict-then-optimize lesson (Elmachtoub and Grigas, 2022; Mandi et al., 2024) and the reason the structure diagnostic of Section 5.4 must be re-run under the decision metric rather than inherited from prediction-metric audits.

Specialists are linear decision-focused predictors per regime (*heads*): head $w \in \mathbb{R}^d$ predicts $\hat{y} = X_t w$ and trains on the SPO+ surrogate (Elmachtoub and Grigas, 2022) of Definition 1’s regret, whose subgradient requires only two calls to the decision oracle per sample. Episode risk is $L_{r,e}(w) = \mathbb{E}_{P_{r,e}}[\rho(Xw; y)]$.

3.2 Capacity and memory models

Each specialist maintains at most $K < R$ regime heads, under two memory models spanning the discrete and graded readings of capacity (Li, 2026): **hard (LRU)**: learning a non-held regime at full capacity evicts the

least-recently-used head (reset to prior), modeling slot-like capacity (an analyst’s active playbooks, a finite adapter store); **soft (interference)**: a shared backbone with per-regime adapters in which every foreign update decays untouched adapters’ effective knowledge by a factor $(1 - \rho_{\text{int}})$, modeling parameter interference in shared representations. Episode buffers (bounded: the most recent E_{buf} episodes, ≤ 40 days each) count toward the same budget and are evicted with their head: knowledge lost is data lost, not merely parameters reset.

3.3 The RISP mechanism

RISP runs N specialists through a batched competition cycle (Algorithm 1):

1. **Shadow evaluation.** Every specialist produces (shadow) decisions for the current regime r_t ; per-specialist decision regret accumulates over a rolling *competition window* of W_c trading days within the regime.
2. **Batched winner-take-all.** At window close, the *window winner* minimizes cumulative window regret plus a niche shaping term $-\lambda(\alpha_{i,r_t} - 1/R)$; only the winner applies one exponentiated-gradient (Hedge) update to its niche affinity $\alpha_i \in \Delta^R$ (Kivinen and Warmuth, 1997; Freund and Schapire, 1997) and trains on the window’s data. Batching is the financial adaptation: per-period regret differences are noise-dominated at realistic signal-to-noise, and a W_c -day cumulative comparison drives the wrong-winner probability down exponentially in W_c (Remark 3); $W_c \in \{1, 5, 20, 60\}$ is ablated in Section 5.11.
3. **Invariant niche training.** A specialist updating regime r minimizes, over its episode buffer for r ,

$$\mathcal{L}_r(w) = \text{Var}_e[\bar{\ell}_{r,e}(w)] + \beta \mathbb{E}_e[\bar{\ell}_{r,e}(w)], \quad (1)$$

where $\bar{\ell}_{r,e}$ is the mean SPO+ surrogate on stored episode e : the regime-indexed Inv-PnCO objective (Yan et al., 2026; Krueger et al., 2021) with episodes as environments. ERM ablations minimize the mean alone.

4. **Pinning.** When a specialist’s affinity for a regime exceeds 0.95, the assignment is *pinned*: the owner alone serves and trains that regime thereafter, making the converged assignment reward-independent—and, because the owner alone trains its niche, isolated—by construction (competition re-opens only if the regime inventory changes).

Regime labeling. All arms (ours and baselines) receive the same causal regime label r_t , computed from information through $t - 1$ only: **L1** (headline), a rolling-percentile volatility band crossed with trend sign; **L2**, a causal two-state filtered volatility model (forward-only Hamilton-style filter (Hamilton, 1989)) crossed with trend sign. The comparison is therefore information-symmetric (task-incremental in continual-learning terms (Li, 2026)): what differs across arms is solely whether the induced capacity assignment chases current reward, not what each is told.

Episode scarcity. Crisis regimes yield 5–7 historical episodes in long equity samples and fewer in shorter panels; $\text{Var}_e[\cdot]$ over so few draws is fragile. Two mitigations are part of the default design: *sub-episode environments* (non-overlapping blocks within episodes enlarge the environment set while preserving between-episode heterogeneity) and *cross-asset episodes* (the same crisis in different asset classes counts as distinct environments, upgrading the claim to cross-asset invariance within a regime). Our synthetic heterogeneity sweep (Section 5.10) controls the episode count directly; the estimation price of few episodes appears explicitly in Lemma 1.

4 Theory: A Decomposition of Post-Reactivation Regret

We state the assumptions, the two lemmas that control each axis, and the propositions that assemble them. Proofs not given inline appear in the companion deep-dive document; all results are idealized-model statements in the register of the GAUSE theory (Li, 2026)—guides to mechanism, with the quantitative claims resting

Algorithm 1 RISP (batched competition cycle)

```
1: Require:  $N$  specialists with capacity  $K$ , regimes  $\{1..R\}$ , window  $W_c$ , EG rate  $\eta$ , niche bonus  $\lambda$ ,  
invariance weight  $\beta$   
2: Initialize  $\alpha_i \leftarrow 1/R$ ; heads empty; no pins  
3: for  $t = 1, 2, \dots, T$  do  
4:   Observe  $X_t$ , regime label  $r_t$  (causal labeler); serve decision of pinned owner if  $r_t$  pinned, else of  
 $\arg \max_i \alpha_{i,r_t}$   
5:   if  $r_t$  pinned then  
6:     Owner stores  $(X_t, y_t)$  in episode buffer; SGD steps on Eq. (1)  
7:   else  
8:     Accumulate window regret  $q_i += \rho_i$  for all  $i$ ; buffer the day  
9:     if window length =  $W_c$  or regime block ends then  
10:       $i^* \leftarrow \arg \max_i [-q_i/|\text{win}| + \lambda(\alpha_{i,r_t} - 1/R)]$   
11:      EG update:  $\alpha_{i^*,r_t} *= e^\eta$ ; renormalize  $\alpha_{i^*}$   
12:       $i^*$  trains on the window's data (Eq. (1)); LRU/interference bookkeeping  
13:      if  $\alpha_{i^*,r_t} \geq 0.95$ : pin  $r_t \rightarrow i^*$   
14:    end if  
15:  end if  
16: end for
```

on the controlled experiments. We do *not* claim adversarial-regret guarantees for the coupled winner-take-all dynamics.

Assumption 1 (Episode exchangeability; regime-indexed invariance). *For each regime r , episode environments are i.i.d. draws $\xi_{r,1}, \xi_{r,2}, \dots \sim Q_r$ from a regime-level hyper-distribution, and the deployed (next) episode is a fresh draw. For a fixed head w define $\mu_r(w) = \mathbb{E}_{\xi \sim Q_r}[L_\xi(w)]$ and $\sigma_r^2(w) = \text{Var}_{\xi \sim Q_r}[L_\xi(w)]$.*

This is the regime-indexed form of the Inv-PnCO environment assumption (Yan et al., 2026): “the next crisis is a new draw from the crisis distribution.” It idealizes on two counts—real episodes are neither independent (macro memory) nor exchangeable (secular drift)—and we flag both in Section 7.

Assumption 2 (Dormancy and foreign updates). *While regime r is dormant for D steps, an allocation mechanism A routes learning updates among specialists. Under a reward-driven A , the specialist holding r is made to learn a currently active regime with per-step probability at least $p > 0$ (the marginal- p form: active regimes are the only sources of reward, reward attracts capacity, and the allocator spreads its updates across the active regimes rather than fixating on one). Call an assignment isolated for r if the owner of r receives no updates other than to r itself during r ’s dormancy—under the hard model this excludes acquiring foreign heads (eviction pressure), and under the soft model it excludes any non- r update, since each one decays r ’s adapter by the interference factor—and covering if every regime has an owner within capacity. (The stronger reading matters: a specialist holding $\{r, r''\}$ with r'' active during r ’s dormancy is isolated in the hard sense but not in the soft sense, and carries a positive soft-model deficit accordingly.) Isolation-plus-coverage is the operative property throughout; reward-independence—an assignment measurable with respect to information independent of the realized reward stream (pinned converged competition, fixed niches, a frozen gate, hand-designed silos)—is the mechanism route by which our arms obtain isolation, not the property itself (Remark 1).*

Remark 1 (Isolation, not reward-independence, does the work). *Reward-independence is neither necessary nor sufficient for isolation: a round-robin trainer is reward-independent and covering yet not isolated—every specialist receives foreign updates on schedule, so its forgetting deficit is positive—while a hand-designed silo is isolated regardless of whether its designer consulted past rewards. The phrase “reward-independent assignment” (Li, 2026) names our mechanism class; the property invoked by the theory below is isolation plus coverage.*

4.1 The invariance axis

Lemma 1 (Episode transfer). *Under Assumption 1, for any fixed w and $\delta \in (0, 1]$, with probability $\geq 1 - \delta$ over the test episode $e' \sim Q_r$,*

$$L_{r,e'}(w) \leq \mu_r(w) + \sigma_r(w)\sqrt{(1-\delta)/\delta} \leq \mu_r(w) + \sigma_r(w)/\sqrt{\delta}.$$

Moreover, with E_r training episodes and $L \in [0, B]$, the empirical mean satisfies $|\hat{\mu}_r(w) - \mu_r(w)| \leq B\sqrt{\log(2/\delta)/(2E_r)}$ with probability $\geq 1 - \delta$ (Hoeffding over episodes), and an analogous $O(B/\sqrt{E_r})$ bound holds for the empirical standard deviation via empirical-Bernstein concentration (Maurer and Pontil, 2009).

Proof. The first display is Cantelli’s one-sided inequality applied to the random variable $L_\xi(w)$, $\xi \sim Q_r$: $\Pr[L_\xi \geq \mu_r + \tau] \leq \sigma_r^2/(\sigma_r^2 + \tau^2)$, solved at the δ level. The second is Hoeffding’s inequality for $[0, B]$ -bounded i.i.d. variables. $\square$

The RISP objective (Eq. 1) is, up to the $O(B/\sqrt{E_r})$ estimation terms, a Lagrangian scalarization of the bound’s two arguments: sweeping β traces the (μ_r, σ_r) Pareto frontier, and for each tail level δ some $\beta(\delta)$ optimizes the bound. ERM ignores σ_r entirely and can sit far up the variance axis—the mechanism by which it loads on episode-specific regularities (Section 5.10 constructs this failure explicitly). We are explicit about register: the trained $\hat{w}$ depends on the training episodes, so applying Lemma 1 to $\hat{w}$ is a plug-in heuristic at the same level as Inv-PnCO’s use of its Theorem 1, not a uniform-convergence guarantee; the estimation terms are the formal price of few episodes.

Lemma 2 (Divergence-to-regret bridge). *For any episode laws P, P' on $\mathcal{X} \times \mathcal{Y}$ and any measurable head w ,*

$$|L_P(w) - L_{P'}(w)| \leq B \cdot \text{TV}(P, P') \leq B\sqrt{\text{KL}(P \| P')}/2.$$

Consequently, if a representation renders episode laws ε -invariant in KL (as in Inv-PnCO-style guarantees, which bound divergences between solution distributions (Yan et al., 2026)), then $\sigma_r(w) \leq B\sqrt{\varepsilon}/2$ for heads factoring through it, and Lemma 1 applies with that variance.

Proof. Regret is $[0, B]$ -valued pointwise (Definition 1), so the difference of its expectations under P, P' is at most $B \cdot \text{TV}$ by the variational characterization of total variation applied to $g = \rho/B \in [0, 1]$; Pinsker’s inequality (Tsybakov, 2009) gives $\text{TV} \leq \sqrt{\text{KL}/2}$. $\square$

This closes the announced technical gap: invariance statements in the divergence metric transfer to the regret metric at the explicit price of the regret range $B = 2k w_{\max} y_{\max}$ (loose—empirical regret ranges are roughly fifty times smaller—but explicit; the decomposition statement below survives verbatim in the KL metric if the constant is objectionable).

4.2 The retention axis

Lemma 3 (Eviction rate under reward-driven allocation). *Hard memory model, capacity K , Assumption 2 with $W_a \geq K$ distinct active regimes during the dormancy of r . Let U_D be the number of distinct non-held foreign regimes the owner of r is made to learn during D dormant steps. Then r ’s head is evicted on the event $\{U_D \geq K\}$, and $\Pr[U_D < K] \leq \Pr[\text{Binom}(D, p(W_a - K + 1)/W_a) < K] \rightarrow 0$ exponentially in D . Under the soft model, the retained strength of r after N_D foreign updates is $(1 - \rho_{\text{int}})^{N_D} \rightarrow 0$.*

Proof. Each dormant step independently delivers a foreign-regime update with probability $\geq p$; conditional on an update, a not-yet-held distinct regime is hit with probability $\geq (W_a - K + 1)/W_a$ while fewer than K distinct regimes have been accumulated. The waiting time to K distinct hits is thus stochastically dominated by a sum of K geometric variables, giving the binomial tail; LRU evicts r once K distinct foreign heads displace it. The soft case is immediate from once-per-update multiplicative decay. $\square$

Proposition 1 (The eviction dichotomy: isolation zeroes the forgetting deficit). *Define the forgetting deficit $\Phi(D; A) = \Pr[r \text{’s head lost at reactivation}]$ (hard) or $1 - \mathbb{E}[(1 - \rho_{\text{int}})^{N_D}]$ (soft). Under Assumption 2: (i) for any reward-driven A in the marginal- p form, $\Phi(D; A) \rightarrow 1$ as $D \rightarrow \infty$ at the rate of Lemma 3—reward-driven allocation destroys isolation; (ii) for any isolated covering A , $\Phi(D; A) = 0$ for all D : the owner of r applies*

zero foreign updates during dormancy, so its head—and its episode buffer—are exactly preserved (*hard*) or decay only with the regime’s own drift (*soft*). RISP’s pinned converged competition is one emergent way to obtain isolation and coverage.

Proof. (i) is Lemma 3. (ii): isolation gives $U_D = 0$ deterministically—no eviction and no interference decay is ever triggered—and coverage guarantees the owner exists and holds r within capacity. $\square$

The two regimes are different *functions* of D —one tends to 1, the other is identically 0—which is what makes the dichotomy experimentally sharp (Section 5.9). Direction (i) carries a disclosed boundary: the proof uses Assumption 2’s marginal- p form, in which the reward-driven allocator spreads across the active regimes; a degenerate allocator that *fixates* on a single active regime never accumulates the K distinct foreign heads that LRU eviction requires and can accidentally retain. Reward-chasing destroys isolation in the typical, spreading case—not by logical necessity—and our reward-driven arms all satisfy the marginal- p form by construction. Coverage is the second conjunct, inherited from GAUSE: random fixed niches are isolated for the regimes they happen to hold but can leave a regime unowned, retaining only what they cover; emergent competition guarantees coverage at $N \geq R$ (Li, 2026), which we re-verify in our setting (every seed of the 20-seed assignment diagnostic converges to a one-per-regime covering assignment; Section 5.2).

4.3 Assembly

Proposition 2 (Post-reactivation regret decomposition). *Under Assumptions 1–2, let regime r reactivate after dormancy D into a fresh episode $e' \sim Q_r$, served by allocation A with the owner’s head $\hat{w}_r$ trained on E_r episodes. Let $\text{CVaR}_\delta^{e'}$ denote the mean of the worst δ -fraction of outcomes over the episode draw e' , and let C_{relearn} denote the probe-window regret excess of a from-prior head over the converged retained benchmark $\hat{w}_r$ (an empirical constant of the learning problem)—so the eviction branch’s regret is the retained benchmark’s regret plus C_{relearn} , not C_{relearn} alone. Then, for any tail level $\delta \in (0, 1]$,*

$$\mathbb{E}_A \left[\text{CVaR}_\delta^{e'} [\rho_{\text{react}}(r)] \right] \leq \underbrace{\mu_r(\hat{w}_r) + \sigma_r(\hat{w}_r) \sqrt{\frac{1-\delta}{\delta}} + O\left(\frac{B}{\sqrt{E_r}}\right)}_{G_{\text{inv}}(\delta)} + \underbrace{\Phi(D; A) C_{\text{relearn}}}_{G_{\text{forget}}},$$

where the outer expectation is over the allocation and eviction randomness and the inner tail functional is over the fresh episode draw. At $\delta = 1$ the statement specializes to the expectation bound $\mathbb{E}[\rho_{\text{react}}(r)] \leq \mu_r(\hat{w}_r) + O(B/\sqrt{E_r}) + G_{\text{forget}}$, in which the σ_r term vanishes.

Proof. Pointwise in both sources of randomness, the accounting inequality is $\rho_{\text{react}}(r) \leq L_{r,e'}(\hat{w}_r) + \mathbf{1}\{\text{evict}\} C_{\text{relearn}}$: on retention the served head is $\hat{w}_r$; on eviction the served head is fresh, and its probe-window regret is the retained benchmark’s regret on the same draw plus C_{relearn} , by the definition of C_{relearn} as an excess. The eviction event is a function of the dormancy-period allocation stream and is independent of e' under Assumption 1. Conditional on that event, bound the first term by the CVaR form of Cantelli’s inequality—for any random variable with mean μ and variance σ^2 , $\text{CVaR}_\delta \leq \mu + \sigma \sqrt{(1-\delta)/\delta}$, the same one-sided moment argument whose quantile version is Lemma 1—plus Lemma 1’s $O(B/\sqrt{E_r})$ estimation price; the second term is constant given the event. Take the outer expectation over the eviction indicator, $\mathbb{E}[\mathbf{1}\{\text{evict}\}] = \Phi(D; A)$. $\square$

The registers of the two terms carry the interpretation: allocation controls the *mean* of the reactivation transient (G_{forget} is an expectation over the eviction randomness and survives at every δ), while the invariance objective controls its *tail across episode draws* (the σ_r term enters only for $\delta < 1$). We state the flip side honestly, because it sharpens rather than weakens the claim: in pure expectation ($\delta = 1$) the σ_r term vanishes and the population-level bound is minimized by the mean-optimal (ERM) head—the variance penalty purchases distributional robustness against the next episode draw, not population-mean improvement, and its empirical mean advantage in E1 operates through the finite-episode estimation term, where ERM loads on episode-specific regularities (the E4 construction).

The tail half of that interpretation was put to its pre-registered test (D5) and *failed*. On the 100-seed synthetic battery the A5–A6 gap does not concentrate in the upper quantiles of per-reactivation regret: the gap-to-median ratios are 1.09 (q75), 0.68 (q90), and 0.43 (CVaR@10%), all Holm-n.s. (min corrected

$p = 0.090$), and the CVaR point estimate leans *opposite* the prediction (excess gap -0.00115 ± 0.00103). Per the lodged branch, the tail interpretation is demoted to an *unconfirmed bound*: the invariance effect appears empirically as a uniform shift across quantiles of the reactivation transient, and the σ_r mechanism, while it bounds the tail, is not established as the operative one. The French-L3 quantile panel leans the predicted way (median gap $+0.00062$ rising to $+0.00264$ at CVaR@10%) but was lodged descriptive-only (27 events) and is not scored (`results/e_cvar_panel.json`).

Proposition 3 (Independent controls; the 2×2 prediction). *G_{forget} depends only on $(A, D, K, \text{memory model})$ and is zeroed by an isolated covering assignment regardless of the training objective; $G_{\text{inv}}(\delta)$ depends only on $(\text{objective}, Q_r, E_r)$ and, at every tail level $\delta < 1$, its σ_r term is reduced by the variance-penalized objective regardless of the assignment. The bound is therefore additive in the two design axes, predicting monotone improvement along both axes of the 2×2 . Moreover the realized interaction is predicted super-additive in one specific cell: under a reward-driven A with $\Phi \approx 1$, the served head at reactivation is fresh regardless of its training objective, so the invariance axis is inert in the reward-driven row (G_{inv} improvements act on a head that is not there to serve).*

Remark 2 (Honest register). *Additivity of the bound does not force additivity of realized regret; we therefore measure the empirical interaction index $I = (A1 - A5) + (A1 - A7) - (A1 - A6)$ and report it with the headline (Section 5.2). The super-additivity prediction—invariance inert in the reward-driven row—is the distinctive falsifiable signature separating our account from a generic “both tricks help” story.*

Proposition 4 (Break-even dormancy). *Charge a carrying cost c per idle specialist-day (infrastructure, attention, calibration operations) and let V scale one unit of probe-window regret to portfolio value over the N_p -day probe. Retention of regime r ’s specialist pays over a dormancy-reactivation cycle iff*

$$c \cdot D \leq \Phi(D; A_{\text{rd}}) C_{\text{relearn}} N_p V,$$

with Φ the reward-driven deficit of Lemma 3. Since Φ saturates at 1 while cD grows linearly, the paying region is a window $[D_{\min}, D_{\max}]$: dormancy long enough that eviction is near-certain, short enough that cumulative carrying cost does not swamp the one-shot reactivation saving. Both endpoints are computed from measured quantities in Section 5.9.

Remark 3 (Batching the competition: a prediction revised by experiment). *The window winner is decided by a W_c -day cumulative regret comparison. For two specialists with true within-regime quality gap Δ and per-day comparison noise σ_n (sub-Gaussian), Hoeffding gives $\Pr[\text{wrong winner in one window}] \leq \exp(-W_c \Delta^2 / 2\sigma_n^2)$, which motivated our pre-registered expectation that per-step competition ($W_c = 1$) would converge to noise at realistic daily signal-to-noise. The experiment revises this: the assignment is decided not by any single window but by the accumulation of EG updates across many windows, and that accumulation performs the averaging the single-window analysis assigns to batching—post-reactivation regret is flat across $W_c \in \{1, 5, 20, 60\}$, with $W_c = 1$ marginally best (Section 5.11). The single-window bound is correct but answers the wrong question; the operative statistic is the win-rate differential integrated by the affinity dynamics. We retain $W_c = 20$ as the default for interpretability (one affinity update per fortnight is auditable) and report the sweep.*

5 Experiments

5.1 Protocol, arms, and metrics

Design philosophy. Every claim is pre-registered as an ordering prediction before the run, tested across 100 seeds with Welch tests and Holm–Bonferroni correction (the controlled battery E1–E6 and the stitched daily panel of Section 5.5; the intraday runs of Sections 5.4 and 5.6 use 20 seeds as run), and reported whether or not it survives—including two predictions that did *not* (the batching expectation, Remark 3, and the real-data dissociation, Section 5.4). All numbers trace to released per-experiment result files. We use three data tiers, in decreasing order of experimental control and increasing order of realism: (i) *synthetic* regime-switching markets with known ground truth (the headline, because only here is every causal claim testable); (ii) *semi-synthetic* regime-stitched real data (real per-regime cross-sections re-arranged into controlled dormancy schedules); (iii) *real walk-forward* panels (five Bybit USDT crypto pairs, daily, 2021–2025; FRED commodities,

2015–2025; and, added in revision, the Ken French industry portfolios, daily—49 industries 1990–2025 plus a withheld-era replication on the 36 complete industries 1926–1989, Section 5.7—portfolio-level equity evidence). We state plainly that the roadmap’s equities flagship (CRSP S&P 500 *constituents*, 1990–2025, with 5–7 crisis episodes) remains the pre-committed forward test: the industry-portfolio results of Section 5.7 are facts about the aggregate cross-section, and aggregation itself can manufacture slow dynamics (Lo and MacKinlay, 1988)—the constituent-level test is where that interpretation is decided.

Synthetic market. $n = 20$ assets, $R = 4$ regimes (three common, one rare “crisis”), $d = 20$ features per asset: two *invariant* dimensions carrying the regime’s decision-relevant structure θ_r (distinct, well-separated across regimes, stable across episodes), one *spurious* dimension whose loading $\gamma_{r,e}$ is redrawn each episode (the episode “weather”), sixteen distractors, and a constant. Signal-to-noise is calibrated to realistic daily cross-sections ($|\theta_r| \approx 0.012$ against noise 0.02–0.06; crisis noise $3\times$ calm), which makes fresh relearning genuinely slow—the regime in which the allocation question is non-trivial, and, as our audit shows (Section 5.12), the regime in which it *stays* non-trivial. Common regimes alternate in blocks of 25–60 days; the crisis regime appears for 15–30 days roughly every fourth cycle, giving dormancy spells of 250–450 days and 15 evaluated reactivations per run ($T \approx 2,370$; first half burn-in, second half evaluated). Decision layer: $k = 4$ of 20, $w_{\max} = 0.25$.

Arms. All learner arms share the identical specialist class (linear SPO+ heads, capacity $K = 2$ of $R = 4$ unless swept, identical learning rates and buffers); they differ *only* in allocation and objective. Capacity accounting: pool arms field $N = 4$ specialists ($4 \times K=2$ heads in total) while monolith arms run a single specialist, so the equal-total-capacity allocation contrast is A5 vs. A2, with the monolith contrasts audited against the E2 capacity sweep (Section 5.8). A1 capacity- K monolith (ERM; always trains the active regime—the canonical reward-driven allocation and the regime-conditioned finance baseline under a capacity budget). A2 mixture-of-experts with a gate trained on realized regret (reward-driven routing). A3 recent-performance capital allocator: capital shares softmax-weight trailing 60-day performance; each specialist trains on the active regime with probability equal to its capital share (capital is recalibration intensity—the industry baseline; its dormant-specialist erosion rate is governed by the softmax floor). A4 random fixed niches (reward-independent, coverage gaps possible). A5 RISP-ERM (emergent assignment, ERM training). A6 RISP-full (emergent assignment, invariant objective, Eq. 1). A7 monolith with the invariant objective (tests whether the objective can substitute for allocation). A8a Hedge over fixed strategies; A8b Hedge over capacity-bounded learning experts (training follows the top weight)—the adaptive online-learning pair of the scope statement. A9 oracle-pinned assignment with ERM; A10 oracle-pinned with the invariant objective (the skyline: hand-given perfect assignment, nothing to discover).

Metrics. *Post-reactivation regret*: mean daily decision regret over the first 15 days of every block whose regime was dormant ≥ 90 days (≥ 80 in the stitched-crypto schedules, whose blocks are shorter; the probe window mirrors the GAUSE protocol). *Overall*: mean daily regret over the evaluation half. *Steady*: within-block regret beyond the probe window (the converged-service level). Where percentages of “excess” regret are quoted, the floor is the steady-state regret of the A10 skyline (0.0252)—the best converged *regime-faithful* service in this market, not an irreducible bound: episode-chasing arms trade *below* it in steady state (A7 0.0243, A8b 0.0244, A1 0.0245) by fitting episode-specific signal, and pay for it at reactivations—the E6 trade-off at baseline SNR. We always quote raw numbers alongside, and we account for this steady-state tax explicitly in Section 5.2.

5.2 E1 — the headline 2×2 dissociation

Table 1 and Figures 1–2 give the headline. Every pre-registered prediction is confirmed (Holm-corrected Welch tests over 100 seeds):

The retention axis is real. Emergent reward-independent assignment with identical ERM training cuts post-reactivation regret from 0.0339 (A2, the regret-trained router) to 0.0324 (A5), -4.4% , $p = 2.5 \times 10^{-7}$ (Holm 1.5×10^{-6}). We headline the router contrast because it is the clean allocation-only comparison: A2 and A5 deploy identical total capacity ($N=4$ specialists of $K=2$ heads each) and differ solely in whether the assignment chases reward. The monolith comparison is confirmatory but not pure: A1 matches the

Table 1: E1: post-reactivation, overall, and steady-state decision regret (mean $\pm$ 95% CI over 100 seeds; synthetic market, $K=2$, hard memory). Horizontal rules group: skylines / RISP / retaining baselines / reward-driven arms / adaptive pair / replay monoliths. Pre-registered ordering confirmed: $A6 < A5 \lesssim A4 < A1 \approx A2 \approx A7 < A3$, with A8a worst overall and A8b tracking the monolith. The A1r replay rows (added in revision, pre-registered D1) are from the replay battery run under identical conventions (`results/e1_replay.json`, 100 seeds), whose re-run of the shared arms reproduces every value shown here to the quoted precision.


Arm	Post-react.	Overall	Steady
A10 oracle assignment + INV (skyline)	0.0306 ± 0.0004	0.0261 ± 0.0002	0.0252
A9 oracle assignment + ERM	0.0322 ± 0.0004	0.0275 ± 0.0003	0.0263
A6 RISP-full (emergent + INV)	0.0308 ± 0.0004	0.0262 ± 0.0002	0.0252
A5 RISP-ERM (emergent + ERM)	0.0324 ± 0.0004	0.0276 ± 0.0003	0.0265
A4 random fixed niches (ERM)	0.0332 ± 0.0005	0.0285 ± 0.0005	0.0274
A1 monolith (ERM)	0.0339 ± 0.0004	0.0268 ± 0.0002	0.0245
A7 monolith (INV)	0.0339 ± 0.0004	0.0266 ± 0.0002	0.0243
A2 MoE router (regret-trained gate)	0.0339 ± 0.0004	0.0271 ± 0.0002	0.0249
A3 recent-performance allocator	0.0375 ± 0.0004	0.0294 ± 0.0002	0.0267
A8a Hedge over fixed strategies	0.0440 ± 0.0006	0.0384 ± 0.0003	0.0372
A8b Hedge over learning experts	0.0338 ± 0.0004	0.0267 ± 0.0002	0.0244
A1r monolith + decoupled replay (ERM)	0.0329 ± 0.0004	0.0277 ± 0.0002	0.0264
A1r ⁺ monolith + replay (INV)	0.0314 ± 0.0004	0.0264 ± 0.0002	0.0252

per-specialist capacity, not the total head count, so its identical gap (0.0339 vs. 0.0324, $p = 1.9 \times 10^{-7}$, Holm 1.3×10^{-6}) partially reflects total heads—E2’s hard monolith at $K=4$ reaches $0.0322 \approx A5$ ’s 0.0324 (Section 5.8). Random fixed niches (A4) trail A5 (0.0332 vs. 0.0324; at 100 seeds the gap reaches significance, $p = 0.011$, Holm 0.044) and carry $1.4\times$ its seed standard deviation—the signature of coverage gaps: a seed whose random draw covers all four regimes retains like A5; a seed with an unowned regime pays for it on every reactivation.

The invariance axis is real—inside a retaining allocation. Holding the emergent assignment fixed, the episode-variance objective cuts post-reactivation regret from 0.0324 (A5) to 0.0308 (A6), -5.0% , $p = 5.8 \times 10^{-9}$ (Holm 4.7×10^{-8}). Jointly, A6 improves on A1 by 9.1% raw ($p = 6.0 \times 10^{-23}$), which is 36% of the excess over the converged invariant-service floor ($0.0339 - 0.0252 = 0.0087$ down to 0.0056).

The steady-state tax, charged in full. The probe-window gain is not free. At steady state the plain monolith trades *below* the invariant arms (A1 0.0245 vs. A6 0.0252, $\approx +0.0007/\text{day}$ against A6) by chasing episode-specific signal—the same behavior E6 shows becomes catastrophic at high SNR, profitable here at baseline. Over this market’s schedule (≈ 225 probe days vs. ≈ 960 steady days in the evaluation half) the two effects largely net out, which is why A6’s *overall* margin over A1 is 2.1% (0.0262 vs. 0.0268), an order smaller than the post-reactivation headline. We state this plainly rather than let the probe metric carry it silently: the mechanism’s value is concentrated at regime onsets—which is where financial decisions are most consequential, and which is the pre-registered reason post-reactivation regret is the primary metric—and a deployment whose loss function weights all days equally should expect the overall number, not the probe number.

The interaction is super-additive, exactly as Proposition 3 predicts. The invariant objective in a reward-driven monolith (A7) is indistinguishable from ERM (A7 vs. A1 post-reactivation: $p = 0.89$; the equivalence also holds positively under a paired TOST with margin ± 0.0008 —pre-declared as half the smallest confirmed effect in this family, the A5-to-A6 invariance gap $|\Delta| = 0.0016$ —at $p = 9.6 \times 10^{-29}$): an objective cannot help a head that was evicted during dormancy. Per-seed, the interaction index $I = (A1-A5) + (A1-A7) - (A1-A6) = -0.0017 \pm 0.0002$ (95% CI excluding zero): the joint improvement exceeds the sum of the marginal improvements because the invariance margin only exists where retention preserves the head it trained. This asymmetry is the distinctive signature of the two-mechanism account—a generic “both tricks help additively” story would put I at zero, and a shared-cause story would make A7 help. (A7 is



(a) The 2×2: the invariance axis is inert in the reward-driven row (the predicted asymmetry).

(b) Daily regret aligned on reactivation.

Figure 1: **Left:** post-reactivation decision regret across allocation × objective (mean ± 95% CI, 100 seeds). Cells: A1 / A7 (top), A5 / A6 (bottom). **Right:** mean daily regret around reactivations (5-day MA, 95% CI): the monolith pays a relearning transient concentrated in the probe window; RISP-full tracks the oracle-assignment skyline from day one.

directionally better than A1 *overall*—0.0266 vs. 0.0268, not significant—consistent with the objective helping while heads are alive and being erased at each eviction.)

Reward-driven arms cluster together, as the theory says they must. The learned router tracks the monolith ($p = 0.88$; TOST equivalence at the same ± 0.0008 margin, $p = 8.2 \times 10^{-12}$); Hedge over learning experts tracks it too (0.0338); the recent-performance allocator is *worst* among learner arms (0.0375)—it not only forgets dormant regimes but trains specialists on a capital-weighted blend of foreign regimes throughout, paying interference even in steady state. The sleeping-experts reading of our problem is answered constructively: A8a (fixed strategies) retains trivially yet is dominated by every learning arm everywhere (0.0440 post-reactivation, 0.0372 steady)—retention without within-niche learning is worth little; A8b (learning experts) learns but forgets—adaptivity of the *combination* does not protect the *experts*. The contribution lives exactly in the gap this pair brackets.

RISP matches the skyline it was never given. A6 is statistically indistinguishable from A10—the hand-pinned oracle assignment with the same invariant objective (0.0308 vs. 0.0306, $p = 0.53$; TOST equivalence at margin ± 0.0008 , $p = 2.6 \times 10^{-12}$)—and significantly better than the oracle assignment with ERM training (A9, $p = 1.5 \times 10^{-6}$, Holm 7.6×10^{-6}). As in GAUSE, the value of competition is not beating hand-assignment given the labels (it cannot); it is *recovering the oracle partition automatically*: in every seed of the 20-seed assignment diagnostic (not rerun at 100 seeds) the population converges to a covering one-per-regime assignment (the three common regimes pin at affinity ≥ 0.95 ; the rare crisis regime reaches 0.67–0.87—short of the pin threshold because crisis windows are scarce—but with the correct unique owner serving throughout, an honest wrinkle we report rather than hide).

Replay does not rescue the monolith (pre-registered D1). The cheapest continual-learning remedy short of a pool is a monolith whose episode buffers are *decoupled* from head eviction (the data survives when the parameters are evicted) plus a generous full burst-refit on the retained buffer at every head recreation—an implementation committed in advance (PREREG_FRENCH49.md D) so that the deficit could not survive as an artifact of our own SGD budget. It closes only 56.3% of the A1–A9 deficit (A1r 0.0329; A1r–A9 = $+0.00073 \pm 0.00025$ paired; paired- t $p = 1.4 \times 10^{-7}$, unpaired Welch $p = 0.011$)—*below* our lodged 60% prediction, which activates the registered anti-branch: eviction-coupling is *stronger* than the theory requires, i.e., the deficit is not merely a data-loss effect. And replay remains far above A6 ($p = 7.0 \times 10^{-14}$). Equipping replay with the invariance objective helps (A1r⁺ 0.0314, beating the ERM-trained oracle assignment A9 at Holm 0.023) but still loses to both invariant pool arms (vs. A6 Holm 0.040; vs. A10 Holm 0.017): retention bought in data space at $O(\text{buffer})$ memory is worth less than retention bought in parameter space, and the assignment—not just the storage—carries value.




Figure 2: E1, all eleven arms (synthetic, $K=2$, hard memory, 100 seeds). Color encodes class: green = emergent reward-independent + invariant; blue = reward-independent (ERM or fixed); orange/red = reward-driven; black = oracle skyline. The dissociation separates classes, not architectures: every reward-chasing arm clusters with the monolith regardless of its machinery.

A fourth, unregistered outcome: reward-coupled serving alone is anti-selective (pre-registered D4). The organizational-carrier probe A3' keeps A3's reward-driven capital and serving verbatim but makes training reward-independent: each specialist trains on its own regime's days regardless of capital share—dormant desks' research stays alive; only deployment is reward-routed. The registration lodged three outcomes (A3' strictly between A3 and A5; $\approx A5$; $\approx A3$); the data chose a fourth we did not register, reported here at full prominence: A3' is strictly *worse* than A3 (0.04580 ± 0.00040 vs. 0.03748 ± 0.00044 post-reactivation, 100 seeds; Holm 1.7×10^{-68}). Audited as no artifact: training intensity is matched (1.000 vs. 1.001 specialist-updates/day) and the serving path is inherited verbatim from A3. The mechanism is identifiable in the served-head telemetry: under A3 the capital leader holds a head for the *active* regime on 94.6% of evaluated days—reward-coupled *training* is self-correcting on the serving side, because whoever holds capital has also been training on the regimes it serves—while under A3' it does on only 39.7%: lagged capital routes serving to the previous regime's owner, which under siloed training holds nothing for the active regime. Reward-coupled serving alone is anti-selective at regime switches. The practitioner translation is one sentence: keeping dormant desks' research alive does not help if capital still routes the serving decision.

5.3 E1m — emergent interference under gradient descent

The controlled results above use an explicit memory model (LRU slots or interference decay), inviting the objection that the forgetting is *coded*, *not emergent*. The pre-registered E-X4 battery (PREREG_FRENCH49.md E) removes the model: a shared trunk ($20 \rightarrow 32 \rightarrow 32$, tanh) with linear heads on the shared representation, the *same* architecture for every arm, plain SGD on the same SPO+ surrogate, and nothing is ever evicted—the monolith is one trunk and one readout, A9 pins per-regime heads sharing the trunk, and A5/A6 run the batched competition over trunk-sharing specialists. Forgetting can arise only as representation interference. It does: $\Gamma_{\text{mlp}} = +0.01009 \pm 0.00037$ (A1–A9 paired, 100 seeds, $p = 1.9 \times 10^{-78}$), and it is reactivation-localized—the probe-window gap is $\approx 4\times$ the steady-state gap—interference *measured*, not injected. The A1-vs-A9 contrast is the classical single-head-versus-multi-head gap of task-incremental continual learning (van de Ven and Tolias, 2019; Farquhar and Gal, 2018), and interference carried by a shared representation is well documented (Ramasesh et al., 2021); what this section adds is not the existence of trunk interference but its *allocation* reading: at fixed architecture, with no eviction machinery and no task identity given, the gap opens under reward-driven assignment, is reactivation-localized in decision regret, and is closed by emergent competition that matches the oracle it was never handed. The paired steady-state gap (+0.0026) reflects




Figure 3: Replay with an eviction-decoupled buffer, both substrates (post-reactivation regret, mean $\pm$ 95% CI). Left: synthetic (100 seeds) — the replay monolith closes 56.3% of the deficit, below the pre-registered 60%, and remains far from A6. Right: French 49 L3 walk-forward (20 seeds) — replay collects the deficit (matching A9) yet the invariance-equipped emergent pool still wins.

the single readout’s conditioning-capacity disadvantage; the reactivation-localized *excess*— $+0.0075 \pm 0.0004$, the probe gap net of the steady gap—is the emergent-forgetting quantity proper. Emergent competition again matches the oracle it was never handed (A6 0.0328 vs. A9 0.0329 post-reactivation), pinning the three common regimes in every seed with the crisis niche uniquely owned but unpinned, as in E1. Two boundaries, at the same prominence as the result. First, an honest null: the invariance axis is *absent* under the MLP (A6 vs. A5, $p = 0.13$)—we do not claim the invariance effect beyond the linear specialists on this evidence. Second, a lodged divergence: at $4\times$ SNR the linear E6 inversion does *not* reproduce— Γ_{mlp} grows $\approx 5\times$ instead of vanishing—with the recorded architectural caveat that the MLP monolith is a single readout that every regime switch must overwrite, not the regime-keyed $K=2$ head store of the linear A1, so the two monoliths are not apples-to-apples and the E6 boundary stands as the *linear* result. What this section establishes is the narrower, load-bearing point: under vanilla gradient descent with no eviction machinery at all, a capacity-bounded reward-driven learner loses dormant-regime service, and a reward-independent assignment protects it.

5.4 E0 — the structure diagnostic on real data, and an honest null

Before any regime-aware mechanism is credited on real data, the diagnostic-first protocol requires evidence that regime structure *exists on the decision metric*: an oracle regime-conditioned predictor (one ridge model per causal regime label, walk-forward refit) must beat the same machinery under block-shuffled labels. GAUSE’s oracle audit found crypto and commodities carry no exploitable regime structure under the *prediction* metric (Li, 2026); since prediction-optimal and decision-optimal differ (the core predict-then-optimize lesson), we pre-registered the re-test under decision regret as a finding either way.

The answer (Table 2) is unambiguous: **no exploitable regime structure on the decision metric** in either domain, under either causal labeler—regime-conditioned regret is statistically identical to block-shuffled-label regret (all $|z| < 0.7$), and both are identical to a pooled unconditioned model. The prediction-metric null of GAUSE *extends* to the decision metric here: the PnO distinction, which in principle could resurrect structure invisible to prediction error, does not do so on these panels. We bound the claim honestly: these are small universes (five crypto pairs, three commodities) at daily frequency with a standard momentum/volatility feature inventory and two labelers; richer features, intraday horizons, or larger cross-sections could yet pass the gate. What the protocol forbids is claiming the mechanism’s benefits where the gate fails—which is exactly what a less guarded evaluation would have produced, since the machinery happily trains and “specializes” on




Figure 4: Emergent interference (EI1m). Left: under a shared MLP trunk with nothing evicted, the A1–A9 gap is $\approx 4\times$ larger in post-reactivation probe windows than at steady state — a forgetting transient, measured rather than injected. Right: the deficit grows with SNR under gradient descent, diverging from the linear E6 inversion (unpaired CIs for the linear series); the single-readout caveat is discussed in the text.

Table 2: E0: structure diagnostic on the decision-regret metric (walk-forward, second-half evaluation; shuffled = block-label permutation, 10 draws). Negative gap = conditioning is (insignificantly) *worse* than shuffled. No domain–labeler pair passes; the real-data dissociation is therefore predicted null, and is (Section 5.5).

Domain	Labeler	Real	Shuffled	Gap	z
Crypto (5 pairs, daily)	L1 vol $\times$ trend	0.01220	0.01202	−1.5%	−0.66
Crypto (5 pairs, daily)	L2 filtered HMM	0.01212	0.01204	−0.7%	−0.17
Commodities (3, daily)	L1 vol $\times$ trend	0.00834	0.00832	−0.3%	−0.14
Commodities (3, daily)	L2 filtered HMM	0.00836	0.00838	+0.2%	+0.16

structureless labels (the GAUSE negative control: organization is not structure).

The intraday retest: the gate can pass. The bound above is not vacuous hedging: re-running the identical protocol at intraday frequency with a richer causal feature inventory (ten features from full OHLCV bars—range-based volatility, volume z -score, reversal, multi-horizon and cross-sectionally ranked momentum) locates a substrate where the gate *passes*. Across a pre-registered $\{1D, 4H, 1H\} \times \{L1, L2\} \times \{\text{bar-native, wall-clock}\}$ screen (ten cells): richer features alone do *not* rescue the daily gate ($|z| < 1.1$)—the binding constraint is the episode inventory, not the features—and every wall-clock-window cell fails, while two bar-native cells (whose faster label dynamics multiply episode counts roughly sevenfold) pass the screen. A pre-registered confirmation with 50 fresh shuffle controls and a split-half stability requirement confirms one: **4-hour crypto under the bar-native L2 labeler** (951 episodes; conditioning beats shuffled labels by 3.15%, $z = 2.32$, both halves positive at +3.4% / +2.9%); the 1-hour cell fails confirmation ($z = 1.76$) and is reported as screen-only. This is the first real-data pass of the structure diagnostic on the decision metric in this line of work, and we state its selection caveats at the same prominence as the pass: it is one confirmed cell from a ten-cell screen (fourteen real-vs-shuffled tests counting the daily battery), the confirmation reuses the same data span as the screen, and the selected cell’s gap estimate is biased upward by selection on the maximum. A disjoint-period replication is the committed next test before any claim is built on this substrate. The pass sets up the sharper question of Section 5.6: does a passing gate suffice for the mechanism to pay?

5.5 E1s — the real-data dissociation is null, as the failed gate predicts

We ran the ten-arm comparison (the arms of Table 1 except the A10 oracle-INV skyline and the later-added A1r replay pair) on *regime-stitched* real crypto data: contiguous same-label day-blocks of the actual cross-sections (real returns, real features), re-arranged into controlled schedules with the rare regime dormant ≈ 200 days (100 seeds). No pair survives Holm correction in the pre-registered test family (min Holm $p = 0.17$; post-reactivation regret 0.0129–0.0136 across arms); the one raw trend—A6 below A1 at uncorrected $p = 0.028$ —we disclose but do not claim, and the family-level conclusion remains null under the pre-registered Holm-corrected test. The numerically lowest arm is A8a—Hedge over *fixed, non-learning* strategies, at 0.0129—which sharpens the null rather than softening it: where the gate fails, a baseline that learns nothing edges out every learning arm, ours included. Under a failed E0 gate this is the *predicted* outcome—there is no per-regime structure for a specialist to retain, so retention cannot pay—and we present it as the protocol working, not as the mechanism failing: the same arms, the same code, and the same statistics that produce the sharp dissociation of Table 1 produce a flat table when the precondition is absent. A method whose evaluation cannot go flat on structureless data should not be trusted when it goes sharp. (Fallback F1 of the pre-registered design: where the gate fails on all available real domains, the headline is the controlled study plus the diagnostic protocol; the equities panel, with its richer episode inventory, is the planned real-data test.)

5.6 E1r — dissociation under a *passed* gate: the second gate

The confirmed intraday pass reverses the falsifiable prediction: with decision-metric structure present, the pre-registered ordering $A6 < A5 \ll A1$ *should* emerge. We ran the ten-arm comparison on regime-stitched 4-hour crypto under exactly the confirmed cell (bar-native L2 labels, the rich feature inventory, the rare vol-up regime dormant ≥ 80 bars, probe 15 bars, 20 seeds). The result is a second null. All pre-registered pairs are flat ($p > 0.57$ throughout; A6 0.0058 ± 0.0002 vs. A1 0.0059 ± 0.0002). We state the record plainly: the pre-registered prediction for this experiment was separation, and it was *refuted*. Two features of the table localize the cause. First, *learning now pays*: A8a (Hedge over fixed, non-learning strategies), the numerically best arm on the structureless daily substrate of Section 5.5, is here the *worst* (0.0061)—consistent with the passed gate. Second, the post-reactivation elevation ($\approx +7\%$ over each arm’s own overall regret) is carried *equally by A9, the hand-pinned oracle that cannot forget by construction*: the measured forgetting deficit $G_{\text{forget}} = A1 - A9$ (post-reactivation) is $+0.000002$, zero to four decimals, against $+0.0017$ in the controlled markets of Table 1. Reactivations are hard here because regime onsets are intrinsically hard, not because anything was forgotten: even with dormancy imposed at ≥ 80 bars by the stitched schedule, a fresh head refits the regime within the 15-bar probe at this bar density and signal level—the E6 boundary (Section 5.12) observed in the wild rather than swept in simulation. Structure exists (gate 1), but there is nothing for retention to harvest (gate 2)—the G_{forget} term of Proposition 2, measured here at zero, explains *retrospectively* why no allocation policy separates. We are explicit about the epistemic status: this is a diagnosis made after the refutation, not a prediction that preceded it—the decomposition supplied the vocabulary, and the honest test of the diagnosis is out-of-sample. That test is the upgraded protocol itself, stated as a falsifiable commitment for every substrate we touch next: *two gates*—structure on the decision metric, then a positive forgetting deficit via the A1-vs-A9 counterfactual, both cheap, both causal, both run before any mechanism is bought—with the pre-registered prediction that the dissociation emerges if and only if both pass. Real data now instantiates three cells of the precondition matrix (no-structure $\Rightarrow$ null, predicted in advance; structure without harvestable forgetting $\Rightarrow$ null, diagnosed after the fact; the controlled study’s structure-with-forgetting $\Rightarrow$ sharp dissociation); the equities cross-section, whose crisis regimes recur on multi-year scales at daily frequency, targets the fourth—and is the first substrate on which the two-gate prediction is exposed to refutation. That exposure happened (Section 5.7), and the protocol did not survive it intact.

5.7 E1f — daily equities (49 industries): the deficit exists, and its sign prices the pool

We took the two-gate commitment to the canonical equity cross-section: the Ken French 49 industry portfolios, daily, value-weighted, 1990–2025 (9,067 days, all 49 industries complete; survivorship handled by construction), built from the 2026-05 CRSP database. Design and predictions were pre-registered before each

run (PREREG_FRENCH49.md): the paper’s exact gate and dissociation protocols, 50 shuffle controls from the start, $k=5/w_{\max}=0.2$ for the wide cross-section, an eight-feature return-based causal inventory, and a labeler family disclosed in full—L1/L2 (the paper’s vol-band and filtered-vol labelers; two cells, pre-registered together) and L3 (an index-drawdown $\times$ trend labeler with genuinely slow crisis dynamics, designed *after* seeing the L1/L2 gate results and pre-registered separately before any L3 statistic was computed; family-wise accounting is over all three cells). In revision, five further batteries were added, each lodged in PREREG_FRENCH49.md before its runs: mechanism probes (X1/X3), transaction costs (D3), a replay baseline (D2), 100-seed reruns of the main batteries (D6), and a withheld-era replication on 1926–1989 (E-F). This section reports the whole record.

Gates. We report gate statistics in the permutation register: z against the 50-control shuffle distribution, whose empirical p cannot fall below $1/51 \approx 0.02$; a normal fit to the control distribution gives the indicative p we quote. (The pre-registration’s auxiliary t -based sub-criterion tests against the shuffle *mean* with an SE denominator; it is register-inflated by $\approx \sqrt{50}$ and non-discriminating—L2 passes it while failing z —so the binding criteria are z , the pooled comparison, and split-half stability, as registered.) L1 passes gate 1 cleanly—conditioning beats the shuffle controls by 1.14% ($z=3.18$, empirical $p \leq 1/51$, normal-fit $p \approx 7 \times 10^{-4}$, both halves positive, conditioned beats pooled)—the first *daily-frequency* structure pass in this program, with no screen-selection caveat (the family is three disclosed cells). L2 fails ($z=0.88$). L3 presents the instructive case: it clears the shuffle criterion (+0.87%, $z=2.24$, normal-fit $p \approx 1.2 \times 10^{-2}$, halves +0.86%/ +0.89%) but fails the pooled sub-criterion by a hair (pooled beats conditioning by 0.10%), so by the pre-registered letter it *fails* gate 1. Gate 2, however, inverts the pattern: on the L1 walk-forward the measured deficit is zero-to-negative ($\hat{\Gamma}_{\text{forget}} = -0.00013 \pm 0.00007$ at 100 seeds—the interval now excludes zero from below), while on the L3 walk-forward it is **positive for the first time on any real substrate in this program**: $\hat{\Gamma}_{\text{forget}} = +0.00089 \pm 0.00011$ (100 seeds, reproducing the 20-seed registration battery; the $\pm$ is seed noise conditional on one history; see caveat (i)). One labeler-anatomy fact matters for interpretation and we state it precisely: the *union* of L3’s crisis regimes is not dormant for long stretches—the 2010, 2011, 2015–16, and 2018 near-miss episodes are all in the crisis inventory, recurring every one to three years—while the specific *crisis-with-uptrend* cell sleeps up to 2,953 trading days. The retained specialists therefore do get occasional live-fire recalibration on near-misses between major crises; the deficit is measured net of those rehearsals, which is the economically relevant quantity.

The threshold sweep. The pre-registered threshold-robustness sweep (PREREG_FRENCH49.md C) then delivered the section’s hardest result, reported here at the same prominence as the positive: **the deficit exists only at the pre-registered 15% cutoff**. At 10% and 12% (denser crisis inventory, shorter dormancies) $\hat{\Gamma}$ is null ($+0.00013 \pm 0.00027$, -0.00010 ± 0.00017) and the ordering dissolves; at 20% (sparser inventory) $\hat{\Gamma} = -0.00045 \pm 0.00026$ and the monolith again beats A6. The 15% value was fixed in the pre-registration before any L3 run, so this is not a tuned peak—but it is a *fragile* one, and the sweep refutes our pre-registered robustness prediction (P8). Two readings, both stated: the fragility is an artifact warning (the practitioner’s reading—the effect lives in one labeler specification on one history), or it is a granularity window exactly parallel to Proposition 4’s dormancy window—too fine a threshold and near-miss rehearsals keep every learner fresh ($\hat{\Gamma} \rightarrow 0$, as the E6 axis predicts), too coarse and the qualifying reactivations become too few to harvest within one history (premium dominates)—with 15% landing, by pre-registered luck or by economics, inside the window. The second reading was post-hoc when first written; the withheld era below was its first discriminating test, and neither simple reading survived it intact.

Dissociations at 100 seeds. The ten-arm walk-forward comparison ran on both labelers (registered at 20 seeds, re-run at 100 under the pre-registered parity battery D6 with outcomes unchanged; 100-seed values quoted throughout), dormancy threshold 90 days. Under L1 ($\hat{\Gamma} < 0$): a significant *inversion*—the monolith beats RISP-full (A1 0.01995 vs. A6 0.02026, Holm 1.2×10^{-26}), the oracle-pinned retainer beats A6 (Holm 8.0×10^{-7})—the invariance premium paid with no insurance claim to collect, the steady-state tax of Section 5.2 surfacing in the probe window itself; on *overall* regret the two are within a hair (0.01849 vs. 0.01848), so the L1 inversion is a probe-window transient, not a net loss—and the costs paragraph below shows it is also a *gross-only* phenomenon. Under L3 ($\hat{\Gamma} > 0$): the **full pre-registered ordering emerges**—A6 $0.02128 \pm 0.00004 < A5$ $0.02165 < A1$ 0.02252 , with the capacity-clean retention contrast (A5




Figure 5: Crisis-window transaction costs, both-pay convention. Left: the L3 deficit $\hat{\Gamma}$ strengthens monotonically with the cost tier (gross point at 100 seeds; net tiers at 20 seeds, annotated). Right: the L1 walk-forward inversion reverses net of 25 bps — the inversion, not the deficit, was the gross-only phenomenon.

vs. the total-capacity-matched router A2) at Holm 1.2×10^{-20} , the monolith contrast at Holm 1.2×10^{-40} (impure in total capacity, as in Section 5.2; no capacity sweep exists on this substrate yet), the invariance axis at Holm 1.6×10^{-15} (A6 vs. A5), the joint effect at Holm 5.2×10^{-65} (-5.5%), and A6 beating the hand-pinned oracle (Holm 1.5×10^{-12})—the synthetic study’s signature at one-half the synthetic effect size, and this time the pool arms also lead on *overall* regret (A6 0.01847, A5 0.01846 vs. A1 0.01858): the probe gain is not repaid in steady state here. One new wrinkle is disclosed rather than absorbed: at 100 seeds the *stitched* counterparts, previously flat in $\hat{\Gamma}$, turn marginally positive under both labelers (L3 $+0.00023 \pm 0.00016$; L1 $+0.00023 \pm 0.00015$) while their arm families stay flat (min raw $p = 0.50$ and 0.34)—the schedule-resampling design now leans the deficit’s way without any arm separating, a favorable drift we flag precisely because it was not predicted.

Mechanism honesty: the deficit is not (yet) forgetting. Two registered probes (E-X1, E-X3) asked *why* the L3 deficit exists, and both returned the answer we did not want. The X1 regression of per-reactivation deficits Γ_i on log cell-dormancy, log union-dormancy, and rehearsal count—with event fixed effects across thresholds, exploiting the same calendar crises appearing at multiple thresholds with different rehearsal histories—finds *no* explanatory covariate (all registered slopes n.s., min $p = 0.44$ pooled, within-event likewise; probe-overlap dilution ruled out at $\leq 4.4\%$). The X3 half-life probe is sharper: at the 15% threshold the A1–A9 gap does *not* decay across the 15-day probe window (tail $\geq$ start; exponential fits $R^2 < 0.2$)—a *level offset*, not a relearning transient; where a true transient exists (20%), it relearns in 2–4 days and nets to zero, exactly the E6 mechanism. We therefore execute the rename the X3 registration committed us to: on real data the measured quantity is a **regime-conditional allocation deficit**—the capacity-bounded, reward-driven learner serves crisis windows persistently worse—and *eviction-forgetting* remains demonstrated only in the controlled settings (the linear E1–E6 and the MLP of Section 5.3). The symbol $\hat{\Gamma}$ and its definition (A1–A9) are unchanged, as is the deployment consequence (the counterfactual prices the pool either way); what changes is the mechanism attribution, and we carry the corrected name through the real-data claims below.

Costs: the deficit survives and strengthens. The pre-registered cost battery (D3; both-pay convention—arm and oracle each charged their own turnover—with the arm-only variant as lodged sensitivity) charges crisis-window effective costs of 25/50/100 bps. The net-of-cost deficit *strengthens monotonically with the cost tier*: $\hat{\Gamma}_{\text{net}} = +0.00126 \pm 0.00020 \rightarrow +0.00161 \pm 0.00021 \rightarrow +0.00230 \pm 0.00024$, with the full four-arm ordering (A6 < A9 < A5 < A1) intact at every tier (A6-vs-A1 Holm 4.6×10^{-19} at 25 bps, strengthening

to 1.3×10^{-29} at 100 bps). The mechanism is the registered one: a reactivated monolith churns its top-5 book hardest exactly when spreads are widest (A1 turnover 1.03/day vs. A6 0.62). The lodged arm-only-pays sensitivity is not merely consistent but *identical*: orderings, $\hat{\Gamma}_{\text{net}}$ values, and p -values match the both-pay convention at every tier, because the oracle’s cost term cancels in the regret difference. Our own lodged pessimism—“the overall-regret margin is expected *not* to survive cost adjustment”—is refuted in the favorable direction. And the L1 inversion **reverses net of costs**: A6 beats A1 at every tier (paired -0.00075 ± 0.00010 at 25 bps, $p = 1.2 \times 10^{-16}$) —the inversion, not the deficit, was the gross-only phenomenon. One cost slice is reported as sensitivity, not re-adjudication: charging costs inside the gate-1 comparison flips L3’s pooled sub-criterion (-0.10% gross $\rightarrow +3.0\%$ net); the gate as registered is a gross-metric criterion and remains a FAIL, and we do not re-open it. The practitioner countermeasure these numbers invite was itself pre-registered (PREREG_FRENCH49.md H) and run: a *trade-banded* monolith, identical to A1 except that the served book updates only when more than b names would change, $b \in \{1, 2\}$ (results/e_french49_banded.json). Banding is free—gross regret is unharmed and net regret improves ($b=2$ vs. A1 net Holm 1.3×10^{-3} on L3)—yet it cannot buy the pool’s advantage: A6 beats both banded monoliths net of 25 bps on L3 (Holm $\leq 7.1 \times 10^{-17}$; the A1–A6 net gap remains 97%/79% intact at $b=1/b=2$) and on L1 (residual $+0.00028 \pm 0.00011$, Holm 2.1×10^{-5})—with the honest disclosure that $b=2$ closes 62% of the L1 net-reversal magnitude, so the reversal sentence above keeps its label with the attenuation reported alongside. The lodged magnitude clause *failed* and is disclosed rather than renarrated: the registered bands cut A1’s turnover only 4.1%/15.7% ($\geq 30\%$ was lodged), because A1’s crisis churn comes from daily re-rankings of more than two names, which defeat narrow bands. The conclusion the battery was built to force is the one it delivers: *selection quality, not just churn, carries the deficit*.

Replay: data retention versus parameter retention. The decoupled-replay monolith of Section 5.2 was run on L3 (D2). Replay collects essentially the entire deficit— $\hat{\Gamma}$ over the replay monolith is $+0.00001 \pm 0.00023$, n.s., $\approx 99\%$ closure against 56% in the controlled market—so on this substrate *data* retention matches *parameter* retention against the ERM oracle, consistent with the allocation-deficit reading above. The registered relabeling branch does not trigger: A6 still beats replay (Holm 0.013) and replay+INV (Holm 4.5×10^{-4}). The pool-superiority claim survives its most dangerous baseline, and the reason is informative: what a buffer cannot buy at any size is the invariance-equipped emergent assignment.

The withheld era: 1926–1989. Everything above shares one 36-year history. The registered E-F battery replicated the byte-identical L3 protocol—frozen 15% specification, same features, arms, tests, and the lodged sign hypothesis; zero new researcher degrees of freedom—on an era the design had never seen: the 36 industries with complete daily history, 1926-07-01 through 1989-12-31 (16,983 days, Depression through the 1987 crash). The design follows the pre-discovery-sample tradition of Linnainmaa and Roberts (2018), whose base rate is the right prior here: most in-sample anomalies attenuate or die pre-sample, and generic pre-sample decay predicts attenuation *everywhere*—not a sign record that scores in both directions with the paying window relocating as crisis density shifts; the out-of-sample discipline is the Goyal–Welch standard (Goyal and Welch, 2008; McLean and Pontiff, 2016). Gate 1 delivers its strongest pass in the program ($z = 23.5$; conditioning beats pooled by $+10.4\%$)—and the frozen 15% cell *inverts* ($\hat{\Gamma} = -0.00024 \pm 0.00008$; A1 beats A6, Holm 2.1×10^{-3}): the strongest structure signal we have ever measured coexists with a pool that prices *negatively*, which is what finally demotes gate 1 to a structure screen. The sign rule, by contrast, scored. Six withheld cells were evaluated (thresholds 10/12/15/20% on the full era; 1926–57 and 1958–89 sub-eras at 15%): **five of six are consistent with the lodged rule, in both directions**, including two positive-deficit cells where the full $A6 < A5 < A1$ ordering emerges out-of-sample—10% ($\hat{\Gamma} = +0.00020 \pm 0.00009$, A6-vs-A1 Holm 1.9×10^{-6}) and 1958–89 ($\hat{\Gamma} = +0.00021 \pm 0.00015$, Holm 2.1×10^{-4})—and two negative-deficit cells where the predicted inversion appears (15% full era; 1926–57, Holm 0.023). The single miss is the 12% cell: $\hat{\Gamma}$ negative-significant in the seed-noise register (-0.00010 ± 0.00007) with a flat arm table—a hit under the weak scoring (“dissociation *only if* $\hat{\Gamma}$ significantly positive”; the weak form is one-directional) and a miss under the strong directional scoring (negative $\Rightarrow$ inversion); we report both scorings rather than choose the flattering one.

Scorecard accounting, stated against ourselves. Three qualifications bound what “five of six” can mean. First, the cells are not independent: the 15% full-era cell is the *union* of the two sub-era cells—its inversion largely restates the 1926–57 inversion—and the four threshold cells share one history’s crises.




Figure 6: The positive-deficit window relocates across eras: $\hat{\Gamma}$ versus drawdown threshold for 1990–2025 (peak at the pre-registered 15%) and the withheld 1926–1989 era (peak at 10%, where the denser crisis inventory shifts the granularity window). Fragility replicates in pattern, not in location.

Alternative countings a reader may prefer: 3/4 (thresholds only), 4/5 (disjoint cells), 5/6 (as published, conservatively scored), 6/6 (weak form). Second, the weak/strong scoring taxonomy was formalized in the verdict addendum *after* these runs: the registration lodged the sign hypothesis but not the hit/miss boundary for a negative-significant $\hat{\Gamma}$ with a flat table. Third, the lodged rule was stated for walk-forward designs, and we decline to hide behind that scoping: at 100 seeds the two *stitched* designs now show positive-significant $\hat{\Gamma}$ ($+0.00023 \pm 0.00015$ and $+0.00023 \pm 0.00016$) with flat arm families (min raw $p \geq 0.34$)—strong-form misses if scored. Counted inclusively the record is 5/8 under the strong form and 8/8 under the weak form (the weak form being one-directional, a positive-significant $\hat{\Gamma}$ with a flat table is a strong-form miss, not a weak-form violation); folding in the four NBER cells of the temporal-custody paragraph below, 8/12 strong and 12/12 weak (with the primary walk-forward’s strong-form status register-dependent, as confessed there). The honest summary: the scorecard’s evidential force is closer to two-to-three effectively independent observations, with both directions represented, than to six independent trials—and its first genuinely temporal-custody test is the lodged forward commitment, not anything in this table. (That commitment’s first cell has since resolved as a both-forms hit; see “The first temporal-custody cell” below.) The sweep pattern is equally instructive: the fragility of the 1990–2025 result *replicates in pattern*—the deficit again lives at essentially one threshold, flanked by flat and inverted neighbors—but the window *relocates* (15% $\rightarrow$ 10%) as the era’s crisis density shifts (at 15% the pre-war era spends 41% of its days in the crisis union, against 13% post-1990). Both simple readings of the original fragility are thereby refuted as stated: the fixed-threshold *artifact* reading is confirmed at the letter for 15%-off-sample yet fails as a full account—the phenomenon reappears off-sample, with the full ordering, twice—and the naive *granularity* reading (multi-year crisis-cell dormancy $\Rightarrow$ positive deficit) fails too: multi-year dormancy is present throughout this era and the sign still varies. What survives both eras is the deficit’s sign as the pricing diagnostic, with the paying window’s *location* an era-dependent function of crisis density rather than of the threshold per se.

What tracked the outcomes. Across the two 1990–2025 walk-forward tests, the sign of $\hat{\Gamma}_{\text{forget}}$ tracked the outcome direction both times; gate 1’s pooled criterion tracked neither, and on the withheld era its strongest pass coincided with an inversion. When the sign rule was first formulated it was *postdictive*—written after seeing both 1990–2025 outcomes, sharing arm A1 with the contrasts it prices (its non-circular content is that




Figure 7: The sign-rule scorecard: $\hat{\Gamma}$ (mean $\pm$ 95% CI) for every scored cell, colored and shaped by the observed outcome (full pre-registered ordering / flat / inversion, classified by the pre-stated rule). Positive deficits coincide with orderings and negative deficits with inversions in both eras; the single miss (12% withheld: negative-significant $\hat{\Gamma}$, flat table) is annotated. Walk-forward intervals are seed noise conditional on each window's single history.

A9 sides with the retaining arms exactly when the deficit is positive)—and we said so. The withheld era was its first out-of-sample exposure, in both directions, and it scored 5 of 6. The rule's standing therefore changes register: no longer merely a lodged hypothesis, not yet a validated deployment rule—the record is one country, one aggregate cross-section, one labeler family, and the miss at 12% is real. The forward test is unchanged and pre-committed: $\hat{\Gamma}$'s sign prices the pool on CRSP constituents, with the registration lodged externally (OSF) before any data access, alongside the NBER-labeler forward test whose specification was lodged before NBER dates were ever joined to a panel.

The first temporal-custody cell, resolved. The NBER forward test ran immediately after its external lodging (registration deposited with the Open Science Framework at 22:47 ET, 2026-07-15; NBER dates first joined to the panel at 22:55 ET), and it is the one cell in this paper whose ex-ante status a third party can certify. The labeler has *no free threshold anywhere*: regime = NBER recession state $\times$ 50-day trend, with the causal primary variant using only announcement-lagged information (a recession enters the labeler on the date the NBER committee *announced* it, verified against the committee's published announcement records). Recession-cell dormancies are finally crisis-scale by construction (up to 2,557 trading days—the 1992→2001 and 2010→2020 gaps). Gate 1 *fails* both variants ($z = 1.74$ primary; pooled edges conditioning)—the third independent dissociation of gate 1 from pool value—while the deficit prices positive: $\hat{\Gamma} = +0.00081 \pm 0.00033$ on the causal walk-forward, strengthening to $+0.00105 \pm 0.00033$ net of 25 bps crisis-window costs (the cost-amplification signature, reproduced ex ante on a labeler the cost battery never saw; turnover 1.04 vs. 0.63/day), with the ordering present exactly as the rule demands ($A6 < A1$ at Holm 1.6×10^{-4} gross; the full $A6 < A5 < A1$ chain at Holm $\leq 3.1 \times 10^{-5}$ net, and $A6$ beating the pinned oracle at 1.7×10^{-3} net). Under the lodged scoring this cell is a **hit under both the weak and the strong forms**. One scoring choice is confessed rather than buried: no registration fixed whether the gross or the net register adjudicates the strong-

form chain—gross, A6-vs-A5 is flat (Holm 0.94) and the full chain holds only net of costs, the economically primary register by the cost battery’s own argument, but an argument that predates this cell rather than this scoring decision. Scored gross-only, the cell is a weak-form hit and a strong-form partial; the amendment lodged for the CRSP test fixes the net register in advance so this ambiguity cannot recur. Disclosures at equal prominence: the primary stitched design is flat ($\hat{\Gamma}$ n.s.)—itself a both-forms hit ($\hat{\Gamma} \approx 0$ with a flat table); the non-causal calendar-dated robustness variant reproduces the walk-forward hit but its stitched cell is a weak-form hit and *strong-form miss* (the same pattern as the withheld 12% cell); and the registration’s LORO and era-blocked slices, initially owed, are now run (`results/e_french49_nber_loro.json`) with a mixed disposition we state in full: the deficit survives every one of the 20 single-event exclusions and all calendar exclusions (drop-all-of-2020: $+0.00042 \pm 0.00026$; drop-2008/09: $+0.00036 \pm 0.00030$)—so it is not single-event-carried—but its mass concentrates in the two deep-dormancy recession-entry reactivations (January 2009, dormancy 1,378 days, and June 2020, dormancy 2,445 days; dropping the former alone cuts $\hat{\Gamma}$ to 41% of headline), and the 2010s era block is *negative*-significant for both contrasts. The cell is therefore event-robust in the strict sense but two-recession-concentrated and not era-consistent—materially weaker than the L3 cell’s robustness profile, and we say so. The interpretable reading, flagged as interpretation: the 2010s block contains no recession-cell reactivations at all (no NBER recession reactivated in that decade’s evaluated window), so it consists entirely of expansion-cell events where retention has nothing to protect and pays premium without claims—concentration in long-dormancy recession entries is arguably the registered phenomenon itself, but era-consistency cannot be claimed. Because the labeler carries no threshold, this result also adjudicates the registered fragility discriminator: the granularity-window reading of the L3 threshold sensitivity gains, and the fixed-cutoff artifact reading loses its remaining support.

Position in the finance and continual-learning literatures. Crisis-conditional strategy behavior is a documented finance regularity—momentum profits condition on market state (Cooper et al., 2004), crash in panic states defined by drawdown-plus-volatility indicators close to our L3 labeler (Daniel and Moskowitz, 2016), and short-horizon predictability concentrates in bad times (Rapach et al., 2010; Henkel et al., 2011)—but that literature conditions a strategy’s *expected return* on the state, whereas Elf measures whether the state’s *dormancy erodes a bounded learner’s parameters*: a retention counterfactual ($\hat{\Gamma}_{\text{forget}}$) that is defined and decision-relevant even where conditional returns are flat, and whose econometric ancestor is the post-break estimation-window question (Pesaran and Timmermann, 2007) rather than state-dependent alpha. The substrate is industry-momentum territory (Moskowitz and Grinblatt, 1999), and the canonical warning that industry-level effects need not survive disaggregation (Grundy and Martin, 2001) is one more reason the CRSP constituent test decides the interpretation. In continual-learning terms, $\hat{\Gamma}_{\text{forget}}$ is a decision-regret backward-transfer measure taken against a multi-head oracle (Lopez-Paz and Ranzato, 2017; Chaudhry et al., 2018); unlike replay remedies that spend memory to suppress forgetting unconditionally (Rolnick et al., 2019; Chaudhry et al., 2019), the two-gate rule asks first whether forgetting is worth buying at all—and the replay-equipped monolith, whose buffer is not coupled to head eviction, has now been run as the natural competing baseline in both settings: it buys most of the deficit and still loses to the pool (this section and Section 5.2). Finally, the persistent, non-decaying form of the deficit is the learner-side analogue of slow-moving capital (Duffie, 2010; Mitchell et al., 2007): there the capital, here the training resource, follows recent reward and therefore arrives at the dislocation late—and the cost battery’s finding that the deficit *strengthens* net of crisis-window costs is that literature’s signature, measured in an allocator’s parameters rather than in prices.

Caveats at the same prominence as the result. (i) Each walk-forward design shares one market history per window; seeds vary arm initialization only (the deterministic A8a arm has a zero-width CI), so its intervals quantify implementation noise *conditional on that history*, not sampling uncertainty across markets. The withheld era raises the count to three windows and six scored cells, but every cell is still a single realized history of one country’s aggregate cross-section. (ii) The stitched counterparts no longer simply support the null: at 100 seeds their $\hat{\Gamma}$ turns marginally positive under both labelers while their arm families stay flat (the wrinkle disclosed above)—still consistent with two readings we cannot separate: the walk-forward precision is single-history optimism, or stitching destroys precisely the long-dormancy, cross-block structure that makes retention pay (the L3 crisis inventory is 11 blocks; stitched schedules cannot reproduce a 2,953-day dormancy by construction). The CRSP panel, which supports walk-forward evaluation





Figure 8: E2: post-reactivation regret vs. per-specialist capacity K ($R = 4$; 100 seeds; left: hard LRU, right: soft interference). Reward-independent arms are flat in K ; reward-driven arms close the gap only as $K \rightarrow R$ under hard slots, and never fully close it under interference.


over multiple disjoint eras, is the discriminating experiment. (iii) L3 was designed after the L1/L2 gate results—with full knowledge of which property (dormancy scale) had killed the first two labelers—and pre-registered before its own runs; under a further Bonferroni $\times 3$ over the labeler family, at 100 seeds every primary dissociation claim survives (all $\leq 4.5 \times 10^{-12}$ corrected, including A6-vs-A9), as do the L1 inversion’s three pairs (headline A1<A6 corrected 3.7×10^{-26}). (iv) Transaction costs are no longer deferred on this substrate: the L3 deficit and the L1 reversal are measured under modeled crisis-window costs (25–100 bps, both-pay). What the cost model does not include—market impact, borrow, capacity—and every other substrate in the paper remain gross; the deflated-Sharpe/SPA battery stays reserved for any alpha claim (Section 5.13). (v) The threshold window is era-dependent: within each era the deficit lives at essentially one threshold, and the paying threshold moved (15% $\rightarrow$ 10%) with crisis density across eras—specification robustness in the naive sense is refuted, and any deployment must estimate the window’s location for its own era. (vi) The evaluated reactivation inventory behind $\hat{\Gamma} > 0$ is 27 events (deterministic across seeds; the evaluated half begins in 2008), right-skewed in contribution—the early-COVID windows, 2025-04, 2015-11, and 2011-11 contribute disproportionately, and about a third of events have individually negative $\hat{\Gamma}_i$. Under leave-one-reactivation-out (`results/e_french49_L3_loro.json`) the deficit survives every single-event exclusion (range $[0.00071, 0.00101]$, all significant), survives excluding the entire 2020 calendar year (0.00080 ± 0.00019 ; the two overlapping COVID probe windows share six days, so the calendar exclusion is the clean test), and is independently positive and significant within each era with a meaningful sample (2010s: 0.00077 ± 0.00022 , $n=14$; 2020s: 0.00094 ± 0.00040 , $n=12$), with the same pattern for the joint A1–A6 contrast. Crisis reactivations contribute more than calm ones—as the mechanism predicts—but no single event, nor all of 2020, accounts for the effect (within the single-history frame of caveat (i), which these analyses cannot lift).

5.8 E2 — capacity sweep: when does the dissociation close?

Sweeping per-specialist capacity $K = 1, \dots, 4$ ($= R$) under both memory models (Figure 8): RISP arms are *flat* in K (A6: 0.0308 at every K ; A5: 0.0324)—a converged one-per-regime assignment needs one slot per specialist, and extra capacity is simply unused. Under the hard model the reward-driven arms improve monotonically with K and approach the retaining class only at $K = R$ (monolith: $0.0339 \rightarrow 0.0322$, matching A5 at $K=4$; router likewise), where eviction becomes unnecessary and the allocation question dissolves. The soft (interference) model inverts the diagnosis instructively: with no eviction, capacity slots stop binding and the reward-driven arms are *flat* in K at ≈ 0.0327 —gentler than hard eviction at tight capacity (0.0339 at $K=1$ and $K=2$) but never closing the gap to A6 (0.0308) at *any* capacity, because graded interference from foreign updates erodes dormant heads no matter how many slots exist. Slack capacity is thus one



(a) Dormancy sweep (E3).



(b) Episode-heterogeneity sweep (E4).

Figure 9: **(a)** Post-reactivation regret vs. dormancy D of the focal regime: the reward-driven arms’ deficit saturates once D exceeds the eviction scale (Lemma 3); levels are not comparable across D (schedules differ), arm gaps within each D are. **(b)** Fixing the retaining allocation and sweeping episode heterogeneity isolates the invariance axis: the ERM–INV gap grows $2.6\times$ across the sweep.

audited boundary of the contribution, with a caveat that depends on the memory physics: with $K \geq R$ and hard isolation, conditioning alone suffices (the regime-switching-finance baseline is vindicated exactly there); with shared representations (the realistic reading for parametric models), interference keeps the population’s structural protection valuable at every capacity.

5.9 E3 — dormancy sweep and the break-even window

Controlled schedules place the focal regime’s reactivations after dormancy $D \in \{21, 63, 126, 252, 504\}$ days (Figure 9a). The retention gap (A1 – A6 post-reactivation) grows from 0.0011 at $D = 21$ to 0.0034–0.0045 for $D \geq 63$ and saturates—the signature of Lemma 3: once dormancy exceeds the eviction scale ($\approx K$ foreign-regime acquisitions, here a few dozen days), the monolith’s head is gone with probability ≈ 1 and longer dormancy adds no further loss, while the RISP owner’s deficit is flat at zero by construction. The recent-performance allocator degrades fastest (its capital floor keeps eroding the dormant specialist throughout). Feeding the measured quantities into Proposition 4 (relearning cost $C_{\text{relearn}} \approx 0.004$ per probe day over $N_p = 15$ days; writing the carrying cost as a fraction f of the per-day probe saving $C_{\text{relearn}}V$, the window condition reduces to $fD \leq \Phi(D)N_p$): the lower endpoint is the eviction scale at which Φ saturates—E3 measures $\Phi(21) \approx 0.27$ and $\Phi(63) \approx 0.86$, so $D_{\min} \approx 60$ days—and the upper endpoint is $D_{\max} = N_p/f$. At $f = 1\%$ retention pays for D between roughly 60 and 1,500 days; at $f = 10\%$ the window shrinks to $D \in [\approx 60, \approx 150]$ days—and at carrying costs above $f^* = N_p \max_D \Phi(D)/D \approx 20\%$ of the probe saving, *no* dormancy length pays. The existence of the *upper* break-even is the practically honest point: idle desks are not free, and a fund whose crisis playbook reactivates once a decade may rationally let it decay if maintaining it costs enough.

5.10 E4 — episode heterogeneity isolates the invariance axis

Holding the allocation fixed (emergent, retaining) and sweeping the episode-heterogeneity scale (the spurious loading’s standard deviation relative to invariant signal) from 0 to $2\times$ baseline (Figure 9b): the ERM specialist degrades steadily ($0.0299 \rightarrow 0.0378$ post-reactivation) while the invariant specialist degrades about 20% slower ($0.0289 \rightarrow 0.0352$; equivalently, ERM degrades 25% faster), so the A5–A6 gap widens monotonically from 0.0010 to 0.0026 ($2.6\times$)—the regime-indexed analogue of the Inv-PnCO heterogeneity result, and direct evidence that the invariance margin is purchased against episode variability specifically (at zero heterogeneity the objectives nearly coincide; the small residual gap is the variance penalty acting as a regularizer against the sixteen distractor features). Reactivation is an out-of-distribution event *within* one’s own regime to precisely




Figure 10: E5 ablations (post-reactivation regret, 100 seeds). **Left:** the invariance objective is U-shaped in β and the variance-only control ($\beta = 0$) fails below retained ERM. **Center:** robustness to the competition window W_c (the revised batching prediction, Remark 3). **Right:** pinning the converged assignment matters.

the degree that episodes differ; retention solves dormancy, and this axis is what remains after.

5.11 E5 — ablations

β sweep, and the variance-only control fails as predicted. Post-reactivation regret is U-shaped in the mean-loss weight: $\beta = 0$ (variance only) gives 0.0357—*worse than retained ERM* (0.0324)—confirming in the decision setting the Inv-PnCO finding that the variance term alone is degenerate (a head can achieve zero episode-variance by being uniformly bad); the plateau $\beta \in [0.25, 4]$ gives 0.0308–0.0318; very large $\beta = 16$ degrades to 0.0351—worse than the retained-ERM arm itself (0.0324), which is the honest reading: β multiplies the effective gradient magnitude, so $\beta = 16$ is not “clean ERM recovered” but an unstably large step on the mean term (A5 remains the correct $\beta \rightarrow \infty$ reference, trained at the proper rate). The objective needs both terms, and is insensitive to β across an order of magnitude.

Competition-window sweep. $W_c \in \{1, 5, 20, 60\}$ days yields 0.0306, 0.0306, 0.0308, 0.0310: flat, with per-step competition marginally *best*—the revision of our pre-registered batching expectation discussed in Remark 3. The mechanism’s robustness to its main scheduling hyperparameter is itself a deployment asset.

Niche bonus and pinning. Competition alone ($\lambda = 0$) reaches 0.0318 versus 0.0308 with the default shaping ($\lambda = 0.3$)—the bonus accelerates but does not cause the assignment, consistent with GAUSE’s λ ablation. Disabling pinning costs 0.0325 versus 0.0308: keeping competition perpetually open re-admits exactly the reward-coupling the mechanism exists to remove, though the EG dynamics’ inertia limits the damage.

5.12 E6 — audit: when retention is worthless

This experiment exists because our first implementation refuted us. At the default parameters of an early prototype (signal-to-noise an order of magnitude above realistic daily cross-sections), *no arm separated from any other*: the monolith relearned an evicted regime within two to five days, the probe window scored mostly converged behavior, and retention was demonstrably worthless. Rather than discard the parameter regime, we promoted it to an audit axis. Sweeping the signal-to-noise multiplier over $\{0.5, 1, 2, 4, 8\}$ (Figure 11):

- The monolith’s relearning penalty (post-reactivation over steady regret) grows from $1.26\times$ to $2.01\times$ with SNR—relative relearning cost rises—but its *absolute* steady level collapses ($0.0295 \rightarrow 0.0165$), because strong signal is refit almost instantly and then exploited deeply.
- The RISP advantage over the monolith is $+6.5\%$ at $0.5\times$, $+9.1\%$ at baseline, $+4.9\%$ at $2\times$ —then *inverts*: -18.8% at $4\times$ and -60.0% at $8\times$. At high SNR the episode-specific loading is itself strongly exploitable, and an ERM monolith that greedily refits each episode beats an invariant specialist that on principle refuses to chase it. The skyline arm inverts identically (A9 tracks A6), confirming this is a property of the *regime*, not of our assignment mechanism.




Figure 11: E6 audit: the relearning penalty (post-reactivation regret relative to own steady state) as the signal-to-noise multiplier sweeps $0.5\times$ – $8\times$. The monolith’s relative penalty grows, but its absolute steady level collapses; the RISP advantage inverts near $2\text{--}3\times$ (see text).

The audit therefore draws the mechanism’s boundary as a measured crossover, near $2\text{--}3\times$ our baseline calibration: RISP pays where signals are weak enough that relearning is slow and episode-chasing does not compensate—which is the empirically realistic regime for daily cross-sectional return prediction—and is counterproductive where within-episode signal is strong, fast to fit, and worth chasing. This is the decision-domain analogue of GAUSE’s stale-specialist result under intra-regime drift: reward-independent retention is an asset under dormancy and a liability under exploitable drift, and the honest deliverable is the location of the boundary.

One scope note belongs here rather than only in the limitations: the inversion documented in this section is established for the linear specialist class; under the shared-trunk gradient-descent replication of Section 5.3 the high-SNR comparison diverges (the deficit grows rather than inverts, with the single-readout caveat discussed there), so the E6 boundary should be read as a property of the linear setting until the neural sweep is architecture-matched.

5.13 Statistical protocol and what we deliberately do not claim

All headline comparisons are Welch tests over 100 independent seeds with Holm–Bonferroni correction within each experiment’s pre-registered pair family. At 100 seeds: the controlled batteries E1–E6, the stitched daily panel, the synthetic replay battery (D1), the MLP battery of Section 5.3 (its SNR sweep at 20), and—per the pre-registered parity battery D6—the French L1/L3 walk-forward and stitched batteries of Section 5.7. At 20 seeds as run: the intraday retest of Section 5.6, the assignment diagnostic, and the French cost (D3), replay (D2), mechanism-probe (X1/X3), and withheld-era (E–F) batteries. In the walk-forward designs of Section 5.7 seeds vary initialization only, so those intervals are conditional on a single history; tables report mean $\pm$ 95% CI. The synthetic and stitched experiments expose every arm to identical data streams per seed (paired design). We do *not* report Sharpe ratios, deflated or otherwise, for the synthetic headline—regret against a known optimum is strictly more informative when ground truth exists—and we make no claim of live trading profitability anywhere: on the panels that fail the gates such a claim would be fitting noise, and on the equities panels (Section 5.7) our numbers are decision-regret at portfolio level—now measured *net* of modeled crisis-window transaction costs for the L3/L1 batteries (both-pay convention, 25–100 bps tiers), gross elsewhere—which is still not a P&L claim. For any future real-data deployment passing the gate, the protocol of record is: deflated Sharpe (Bailey and López de Prado, 2014) for any P&L claim, an SPA/Reality-Check test (Hansen, 2005; White, 2000) across the arm family, transaction costs at 5/10/25 bps with borrow, nested walk-forward for every hyperparameter, and the break-even-dormancy carrying-cost analysis of Proposition 4—pre-committed here.

5.14 Synthesis

Across E0–E6 the two-axis account survives every controlled test we constructed against it, and E1m closes its last controlled loophole: the deficit emerges under vanilla gradient descent with nothing evicted. On real data its *protocol* was refuted twice and revised twice (E1r, E1f)—each time toward the decomposition’s own central term. What survives that record, stated as the four claims we would defend: (1) the *sign* of $\hat{\Gamma}$ prices the pool—two in-sample tests and five of six withheld cells, in both directions; (2) where the deficit is positive it is real economics—amplified, not killed, by crisis-window transaction costs, and collectible by replay on real data though the invariance-equipped pool still beats replay on both substrates; (3) its mechanism on real data is honestly open—the measured quantity behaves as a regime-conditional *allocation* deficit, while eviction-forgetting proper remains a controlled-setting result (E1, E1m); (4) gate 1 does not price the pool and is demoted to a structure screen. Retention is an allocation property: every isolated covering arm retains (A4–A6, A9, A10), every reward-chasing arm relearns (A1–A3, A8b), the boundary is crisp ($p \sim 10^{-23}$ for the class separation), and no training objective crosses it (A7). Invariance is an objective property: it pays only inside a retaining allocation (super-additive interaction, CI excluding zero), grows with episode heterogeneity (E4), requires both of its terms (β sweep)—and is honestly bounded: it is absent under the shared-trunk MLP (E1m). The mechanism needs its preconditions: bounded capacity (E2), real dormancy (E3), weak-signal relearning (E6), and verifiable regime structure (E0/E1s)—each audited with the experiment that locates its boundary. Where all preconditions held, the emergent system matched the hand-designed skyline without being given the assignment; where any failed, the system honestly bought nothing.

6 Discussion

Why two mechanisms, not one. The temptation in both literatures is monism. Continual-learning work treats non-stationarity as forgetting and reaches for retention; OOD-generalization work treats it as shift and reaches for invariance. The financial setting forces the synthesis because its non-stationarity is genuinely two-layered: regimes *recur* (so there is something worth retaining) and recurrences *differ* (so retention alone retains the wrong thing). The super-additive interaction is the empirical signature that the layers are real: each mechanism’s value is contingent on the other having done its job. We conjecture this structure—retain the niche, generalize within it—is generic to any domain with recurring-but-drifting modes: climate regimes, epidemic waves, demand seasons, adversary playbooks.

The allocator, not the org chart. Multi-strategy funds already silo desks by regime—the practitioner instantiation of reward-independent assignment. RISP’s contribution to that practice is not the silo but the analysis of what destroys it: recent-performance capital allocation *across* silos (our A3, the industry’s default) erodes dormant desks at a rate set by its capital floor, and a learned router—the modern ML upgrade—fails identically and for the same structural reason (Proposition 1). The emergent mechanism additionally shows the silo map need not be designed: at $N \gtrsim R$ it is recovered automatically, matching the hand-pinned skyline, with graceful degradation documented when competition’s preconditions weaken (the unpinned crisis niche of Section 5.2). The D4 probe sharpens the carrier attribution: keeping every silo’s research alive while letting recent-performance capital route the serving decision (A3’, Section 5.2) is strictly worse than the fully reward-coupled default—the deployment coupling, not training starvation, is the destructive half.

What the diagnostic-first protocol buys. The most consequential result in this paper for applied work is the pairing of a null and its mirror: where the diagnostics refuse to certify structure or forgetting, the machine run anyway goes flat or inverts (E1s, E1r, E1f-L1, the withheld 15% cell)—and wherever the measured deficit is positive, the same machine, unchanged, produces the pre-registered ordering (E1f-L3 in-sample; the withheld 10% and 1958–89 cells out-of-sample). The protocol converts “our regime-aware method beat baselines on a backtest” (unfalsifiable in the presence of label noise and selection) into separable, individually falsifiable claims: *structure exists here* (E0, cheap, runs before any method), *dormancy actually destroys something* ($\hat{\Gamma}_{\text{forget}}$, one counterfactual arm), and *given both, the mechanism captures it* (E1). We commend the split to the regime-aware finance literature at large.

6.1 Claims ledger

Every headline claim, its evidence grade, and what would refute it, in one place; rows ordered strongest to most contingent. Rows 1–6 are settled within their stated scope; rows 7–10 are the frontier; rows 11–12 are the lodged bets.

Claim	Evidence (status)	What would refute it
1. Two orthogonal axes; 2×2 dissociation	100 seeds, joint $p=6.0\times 10^{-23}$; TOST'd equivalences (<i>established, controlled</i> ; D5 quantile test failed—the bound's tail reading demoted to unconfirmed, disclosed)	a shared-backbone setting collapsing one axis
2. Super-additivity: invariance inert in reward-driven allocation	I CI excludes 0; $A7\approx A1$ TOST $p=9.6\times 10^{-29}$ (<i>established, controlled</i>)	A7 helping over A1 anywhere
3. Emergent assignment recovers the oracle map unseen	$A6\approx A10$ $p=0.53$; covering 20/20 seeds (<i>established</i> ; crisis niche unpinned, disclosed)	non-coverage at $N \gtrsim R$
4. Interference emerges under gradient descent, nothing evicted	Γ_{mlp} excess $+0.0075 \pm 0.0004$, $p=1.9\times 10^{-78}$ (<i>established</i> ; invariance null under MLP, disclosed)	the gap vanishing at matched architecture
5. Retention inverts at high SNR (linear class)	E6 sweep; oracle inverts identically (<i>established, linear</i> ; MLP diverges, disclosed)	inversion reproducing nowhere outside E6
6. Structure screen is demoted: it does not price the pool	three dissociations incl. $z=23.5$ pass with inversion (<i>established by record</i>)	the screen tracking pool value out-of-sample
7. The deficit is real economics where positive	cost-amplified $+0.00126 \rightarrow +0.00230$ net; replay collects it, pool still wins; banded-monolith countermeasure run, net edge stands (<i>frontier: one country, aggregate</i>)	a disciplined implementation beyond $b\leq 2$ bands closing the net gap
8. Real-data mechanism = regime-conditional allocation deficit	X1/X3: no dormancy/rehearsal covariate; non-decaying (<i>honestly open</i>)	a covariate or transient found; or CRSP contradiction
9. $\hat{\Gamma}$'s sign prices the pool	8/12 strong, 12/12 weak (one-directional) incl. one temporal-custody hit; dependence caveats stated (<i>frontier</i>)	one clean refuting cell (definitions lodged)
10. Window location tracks crisis density	two eras: 15% $\leftrightarrow$ 10% relocation ($n=2$; <i>estimator lodged for CRSP</i>)	estimator missing on regionals/CRSP
11. Live measurement-split prediction	frozen 2026-06-30, both cells positive; scored 2027-04-07 (<i>lodged, unresolved</i>)	A1 outperforming A6 in the outcome window
12. CRSP T-split forward battery	OSF-lodged design, register and estimator fixed pre-data (<i>lodged, unresolved</i>)	sign-rule miss under the T-split

Table 3: The claims ledger. Grades and refutation conditions are lodged in the pre-registration record (osf.io/nsx4e; [PROVENANCE.md](#)).

7 Limitations

- The controlled study carries the causal claims; the real-data record is multi-era but still one country and one aggregate cross-section.** The dissociation is demonstrated with full causal control only in synthetic markets. On real data, the record now spans 1926–2025 (three walk-forward windows, six withheld scored cells, costs measured on the 1990–2025 batteries), but every cell is portfolio-level, single-history, and drawn from one country's industry aggregate; the crypto/commodity nulls and the L1/prewar inversions bound it, and the mechanism attribution on real data is open—allocation deficit versus eviction-forgetting (Section 5.7)—with CRSP constituents and an NBER-anchored slow labeler as the pre-committed deciders. The constituent-level demonstration remains undone: the CRSP S&P 500 panel (1990–2025, 2000/2008/2011/2015/2018/2020/2022 as crisis episodes) is the pre-committed forward test with its statistics fixed in advance (Section 5.13), and it is where the sign rule, the aggregation caveat, and the mechanism question are decided.
- Real universes remain modest.** Five crypto pairs, three commodities, and 49 (post-1990) or 36 (pre-1990) industry portfolios at daily frequency; the E0/E1s nulls are bounded to their panels, labelers, and feature inventories, and the E1f positives are bounded to one aggregate cross-section under one drawdown labeler family (threshold sweeps pre-registered in both eras).
- Assumption 1 idealizes episodes.** Real episodes are neither independent (macro memory links crises) nor exchangeable (secular drift in market microstructure), and crisis regimes offer 5–7 episodes at

best—the $O(B/\sqrt{E_r})$ estimation terms in Lemma 1 are not decorative. The sub-episode and cross-asset mitigations of Section 3.3 enlarge the environment set but introduce their own dependence; we flag both rather than resolve them.

4. **Theory register.** Propositions 1–3 are idealized-model results: the reward-driven eviction rate assumes per-step foreign-update probability $p > 0$ (verified by construction in our arms, plausible but unverified for arbitrary allocators), and we claim no adversarial-regret guarantee for the coupled winner-take-all dynamics. The bound’s additivity motivates but does not entail the empirical super-additivity; we measure the latter.
5. **The mechanism inverts under strong exploitable drift.** E6 is a limitation stated as an experiment: at high within-regime SNR, invariant retention loses to episode-chasing ERM by up to 60%. Any deployment must locate itself on the E6 axis first; a staleness trigger in the GAUSE style (re-opening competition on confident underperformance) is the known remedy and is future work here. The boundary is established for the linear, regime-keyed model class: under the single-readout MLP of Section 5.3 the high-SNR inversion does not reproduce (the lodged divergence, with its architectural caveat).
6. **Capacity as a binding constraint is a modeling choice.** As in GAUSE, a purely tabular learner could hold all regimes; the contribution is conditional on capacity binding (true for fixed-size parametric models and for the operational reading—calibration attention, infrastructure—of strategy pools). The soft-interference results (E2) show the dissociation does not depend on the discrete eviction model.
7. **Task-incremental scope.** Every arm receives the same causal regime label; the comparison is information-symmetric, but inferring regimes while specializing (class-incremental) is unsolved here. GAUSE’s label-free variant collapsed under real feature overlap, and we expect the financial case to be at least as hard; we do not claim otherwise.

8 Conclusion

We decomposed non-stationary financial decision error into two independently controlled failure modes—between-regime forgetting, owned by the capacity allocator, and within-regime drift, owned by the training objective—and showed theoretically and experimentally that fixing either without the other forfeits most of the available improvement, while fixing both recovers the hand-designed skyline without hand design. The enabling property, imported from GAUSE and ported to decision-focused learning, is the isolation-with-coverage of the converged assignment—obtained emergently through reward-independent pinning; the new axis, imported from invariant predict-and-optimize and ported to recurring regimes, is episode-variance penalization of a decision-focused loss; the bridge between their guarantee languages is one Pinsker step. Around the mechanism we built the evaluation standard we believe the area needs: pre-registered orderings, a structure diagnostic that runs before claims, audits that locate every boundary of validity, and the discipline of reporting the predictions the data refused. When the regime returns, the specialist who kept the playbook—and kept it general—decides best; this paper says precisely when that sentence is true, and how to check.

References

- Rahaf Aljundi, Punarjay Chakravarty, and Tinne Tuytelaars. Expert Gate: Lifelong learning with a network of experts. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2017.
- Andrew Ang and Geert Bekaert. International asset allocation with regime shifts. *The Review of Financial Studies*, 15(4):1137–1187, 2002.
- Andrew Ang and Allan Timmermann. Regime changes and financial markets. *Annual Review of Financial Economics*, 4(1):313–337, 2012.
- Martin Arjovsky, Léon Bottou, Ishaan Gulrajani, and David Lopez-Paz. Invariant risk minimization. In *arXiv preprint arXiv:1907.02893*, 2019.

- Sanjeev Arora, Elad Hazan, and Satyen Kale. The multiplicative weights update method: A meta-algorithm and applications. *Theory of Computing*, 8(1):121–164, 2012.
- David H Bailey and Marcos López de Prado. The deflated sharpe ratio: correcting for selection bias, backtest overfitting, and non-normality. *The Journal of Portfolio Management*, 40(5):94–107, 2014.
- Avrim Blum and Yishay Mansour. From external to internal regret. *Journal of Machine Learning Research*, 8:1307–1324, 2007.
- Olivier Bousquet and Manfred K Warmuth. Tracking a small set of experts by mixing past posteriors. *Journal of Machine Learning Research*, 3:363–396, 2002.
- Aymeric Capitaine, Maxime Haddouche, Eric Moulines, Michael I. Jordan, Etienne Boursier, and Alain Durmus. Online decision-focused learning. In *International Conference on Learning Representations (ICLR)*, 2026. arXiv:2505.13564.
- Arslan Chaudhry, Puneet K. Dokania, Thalaiyasingam Ajanthan, and Philip H. S. Torr. Riemannian walk for incremental learning: Understanding forgetting and intransigence. In *European Conference on Computer Vision (ECCV)*, 2018.
- Arslan Chaudhry, Marcus Rohrbach, Mohamed Elhoseiny, Thalaiyasingam Ajanthan, Puneet K. Dokania, Philip H. S. Torr, and Marc’Aurelio Ranzato. On tiny episodic memories in continual learning. *arXiv preprint arXiv:1902.10486*, 2019.
- Michael J. Cooper, Roberto C. Gutierrez, and Allaudeen Hameed. Market states and momentum. *Journal of Finance*, 59(3):1345–1365, 2004.
- Thomas M Cover. Universal portfolios. *Mathematical Finance*, 1(1):1–29, 1991.
- Kent Daniel and Tobias J. Moskowitz. Momentum crashes. *Journal of Financial Economics*, 122(2):221–247, 2016.
- Amit Daniely, Alon Gonen, and Shai Shalev-Shwartz. Strongly adaptive online learning. *International Conference on Machine Learning*, pages 1405–1411, 2015.
- Darrell Duffie. Presidential address: Asset price dynamics with slow-moving capital. *Journal of Finance*, 65(4):1237–1267, 2010.
- Graham Elliott and Allan Timmermann. Optimal forecast combination under regime switching. *International Economic Review*, 46(4):1081–1102, 2005.
- Adam N Elmachtoub and Paul Grigas. Smart “predict, then optimize”. *Management Science*, 68(1):9–26, 2022.
- Sebastian Farquhar and Yarin Gal. Towards robust evaluations of continual learning. *arXiv preprint arXiv:1805.09733*, 2018.
- Chrisantha Fernando, Dylan Banarse, Charles Blundell, Yori Zwols, David Ha, Andrei A Rusu, Alexander Pritzel, and Daan Wierstra. PathNet: Evolution channels gradient descent in super neural networks. *arXiv preprint arXiv:1701.08734*, 2017.
- Yoav Freund and Robert E. Schapire. A decision-theoretic generalization of on-line learning and an application to boosting. *Journal of Computer and System Sciences*, 55(1):119–139, 1997.
- Yoav Freund, Robert E Schapire, Yoram Singer, and Manfred K Warmuth. Using and combining predictors that specialize. *Proceedings of the 29th Annual ACM Symposium on Theory of Computing*, pages 334–343, 1997.
- Amit Goyal and Ivo Welch. A comprehensive look at the empirical performance of equity premium prediction. *Review of Financial Studies*, 21(4):1455–1508, 2008.

- Bruce D. Grundy and J. Spencer Martin. Understanding the nature of the risks and the source of the rewards to momentum investing. *Review of Financial Studies*, 14(1):29–78, 2001.
- Massimo Guidolin and Allan Timmermann. Asset allocation under multivariate regime switching. *Journal of Economic Dynamics and Control*, 31(11):3503–3544, 2007.
- Zirui Guo, Quan Cheng, Da-Wei Zhou, and Lijun Zhang. Stable routing for mixture-of-experts in class-incremental learning. *arXiv preprint arXiv:2605.17571*, 2026. StaR-MoE.
- James D Hamilton. A new approach to the economic analysis of nonstationary time series and the business cycle. *Econometrica*, 57(2):357–384, 1989.
- Peter R Hansen. A test for superior predictive ability. *Journal of Business & Economic Statistics*, 23(4):365–380, 2005.
- Campbell R Harvey and Yan Liu. Backtesting. *The Journal of Portfolio Management*, 42(1):13–28, 2016.
- Elad Hazan and Comandur Seshadhri. Efficient learning algorithms for changing environments. In *International Conference on Machine Learning*, pages 393–400, 2009.
- Sam James Henkel, J. Spencer Martin, and Federico Nardari. Time-varying short-horizon predictability. *Journal of Financial Economics*, 99(3):560–580, 2011.
- Mark Herbster and Manfred K Warmuth. Tracking the best expert. *Machine Learning*, 32(2):151–178, 1998.
- James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwinska, et al. Overcoming catastrophic forgetting in neural networks. *Proceedings of the National Academy of Sciences*, 114(13):3521–3526, 2017.
- Jyrki Kivinen and Manfred K. Warmuth. Exponentiated gradient versus gradient descent for linear predictors. *Information and Computation*, 132(1):1–63, 1997.
- David Krueger, Ethan Caballero, Joern-Henrik Jacobsen, Amy Zhang, Jonathan Binas, Dinghuai Zhang, Remi Le Priol, and Aaron Courville. Out-of-distribution generalization via risk extrapolation (rex). In *International Conference on Machine Learning*, pages 5815–5826, 2021.
- Saurabh Kumar, Henrik Marklund, Ashish Rao, Yifan Zhu, Hong Jun Jeon, Yueyang Liu, and Benjamin Van Roy. Continual learning as computationally constrained reinforcement learning. *Foundations and Trends in Machine Learning*, 18(5):913–1053, 2025. arXiv:2307.04345.
- Bin Li and Steven CH Hoi. Online portfolio selection: A survey. *ACM Computing Surveys*, 46(3):1–36, 2014.
- Hongbo Li, Sen Lin, Lingjie Duan, Yingbin Liang, and Ness B. Shroff. Theory on mixture-of-experts in continual learning. In *International Conference on Learning Representations (ICLR)*, 2025. arXiv:2406.16437.
- Yuhao Li. Emergent specialization in learner populations: Reward-independent capacity assignment as a defense against catastrophic forgetting, 2026. GAUSE; companion paper and code: <https://github.com/HowardLiYH/GAUSE>.
- Juhani T. Linnainmaa and Michael R. Roberts. The history of the cross-section of stock returns. *Review of Financial Studies*, 31(7):2606–2649, 2018.
- Yang Liu, Toan Nguyen, and Flora D Salim. CP-MoE: Consistency-preserving mixture-of-experts for continual learning. *arXiv preprint arXiv:2605.20247*, 2026.
- Andrew W. Lo and A. Craig MacKinlay. Stock market prices do not follow random walks: Evidence from a simple specification test. *Review of Financial Studies*, 1(1):41–66, 1988.
- Marcos López de Prado. *Advances in Financial Machine Learning*. John Wiley & Sons, 2018.
- David Lopez-Paz and Marc’Aurelio Ranzato. Gradient episodic memory for continual learning. In *Advances in Neural Information Processing Systems 30*, 2017.

- Arun Mallya and Svetlana Lazebnik. PackNet: Adding multiple tasks to a single network by iterative pruning. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 7765–7773, 2018.
- Jayanta Mandi, Emir Demirović, Peter J Stuckey, and Tias Guns. Smart predict-and-optimize for hard combinatorial optimization problems. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 34, pages 1603–1610, 2020.
- Jayanta Mandi, James Kotary, Senne Berden, Maxime Mulamba, Victor Bucarey, Tias Guns, and Ferdinando Fioretto. Decision-focused learning: Foundations, state of the art, benchmark and future opportunities. *Journal of Artificial Intelligence Research*, 80:1623–1701, 2024.
- Andreas Maurer and Massimiliano Pontil. Empirical bernstein bounds and sample variance penalization. In *Conference on Learning Theory*, 2009.
- R David McLean and Jeffrey Pontiff. Does academic research destroy stock return predictability? *The Journal of Finance*, 71(1):5–32, 2016.
- Mark Mitchell, Lasse Heje Pedersen, and Todd Pulvino. Slow moving capital. *American Economic Review*, 97(2):215–220, 2007.
- Tobias J. Moskowitz and Mark Grinblatt. Do industries explain momentum? *Journal of Finance*, 54(4): 1249–1290, 1999.
- Peter Nystrup, Henrik Madsen, and Erik Lindström. Dynamic portfolio optimization across hidden market regimes. *Quantitative Finance*, 18(1):83–95, 2018.
- Peter Nystrup, Petter N Kolm, and Erik Lindström. Greedy online classification of persistent market states using realized intraday volatility features. *The Journal of Financial Data Science*, 2(3):25–39, 2020.
- Chaofan Pan, Lingfei Ren, Linbo Xiong, Yonghao Li, Wei Wei, and Xin Yang. Recap: Regime-adaptive continual learning for portfolio management. In *Proceedings of the 32nd ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD)*, 2026. arXiv:2606.00143.
- German I Parisi, Ronald Kemker, Jose L Part, Christopher Kanan, and Stefan Wermter. Continual lifelong learning with neural networks: A review. *Neural Networks*, 113:54–71, 2019.
- M. Hashem Pesaran and Allan Timmermann. Selection of estimation window in the presence of breaks. *Journal of Econometrics*, 137(1):134–161, 2007.
- Daniel Philps, Tillman Weyde, Artur d’Avila Garcez, and Roy Batchelor. Continual learning augmented investment decisions. *arXiv preprint arXiv:1812.02340*, 2018. NeurIPS 2018 Workshop on Challenges and Opportunities for AI in Financial Services.
- Vinay V. Ramasesh, Ethan Dyer, and Maithra Raghu. Anatomy of catastrophic forgetting: Hidden representations and task semantics. In *International Conference on Learning Representations (ICLR)*, 2021.
- David E. Rapach, Jack K. Strauss, and Guofu Zhou. Out-of-sample equity premium prediction: Combination forecasts and links to the real economy. *Review of Financial Studies*, 23(2):821–862, 2010.
- David Rolnick, Arun Ahuja, Jonathan Schwarz, Timothy P. Lillicrap, and Greg Wayne. Experience replay for continual learning. In *Advances in Neural Information Processing Systems 32*, 2019.
- Andrei A Rusu, Neil C Rabinowitz, Guillaume Desjardins, Hubert Soyer, James Kirkpatrick, Koray Kavukcuoglu, Razvan Pascanu, and Raia Hadsell. Progressive neural networks. *arXiv preprint arXiv:1606.04671*, 2016.
- Joan Serra, Dídac Surís, Marius Miron, and Alexandros Karatzoglou. Overcoming catastrophic forgetting with hard attention to the task. In *International Conference on Machine Learning*. PMLR, 2018.

- Sanket Shah, Kai Wang, Bryan Wilder, Andrew Perrault, and Milind Tambe. Decision-focused learning without decision-making: Learning locally optimized decision losses. In *Advances in Neural Information Processing Systems*, volume 35, pages 1320–1332, 2022.
- Alexandre B Tsybakov. *Introduction to Nonparametric Estimation*. Springer, 2009. Pinsker’s inequality, ch. 2.
- Gido M. van de Ven and Andreas S. Tolias. Three scenarios for continual learning. *arXiv preprint arXiv:1904.07734*, 2019.
- Halbert White. A reality check for data snooping. *Econometrica*, 68(5):1097–1126, 2000.
- Mitchell Wortsman, Vivek Ramanujan, Rosanne Liu, Aniruddha Kembhavi, Mohammad Rastegari, Jason Yosinski, and Ali Farhadi. Supermasks in superposition. In *Advances in Neural Information Processing Systems*, volume 33, 2020.
- Rui Yan, Haoyu Geng, others, Yuhao Li, et al. Invariant predict-and-optimize under distribution shift. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2026. Inv-PnCO; to appear.